ResearchGate

See discussions, stats, and author profiles for this publication at: https://www.researchgate.net/publication/387490884

Apriori Algorithm for Engineers

Book · December 2024

DOI: 10.5281/zenodo.14566774

CITATION
1

READS
55

1 author:

Partha Majumdar
Swiss School Of Business and Management Geneva

171 PUBLICATIONS **9** CITATIONS

SEE PROFILE

All content following this page was uploaded by Partha Majumdar on 28 December 2024.

The user has requested enhancement of the downloaded file.

Apriori Algorithm for Engineers

A FOUNDATION FOR RECOMMENDATION SYSTEMS

Partha Majumdar

Copyright © 2024 Partha Majumdar

All rights reserved.

No part of this book may be reproduced, stored in a retrieval system, or transmitted in any form or by any means, electronic, mechanical, photocopying, recording, or otherwise, without express written permission of the author.

ISBN-13: 9798305158908

Cover design by Partha Majumdar.

Unless explicitly stated, all images are created by the author or licensed from Adobe.

Dedicated to

Prof. (Dr.) Shashi Jain.

Your lectures opened my eyes to this fascinating domain.

Preface

In modern engineering, data-driven decisions are key, and recommendation systems are vital engines in industries and sectors, from retail to healthcare. "**Apriori Algorithm for Engineers**" is a guide for engineers, analysts, and technologists who want to bridge the gap between theory and real-world application and build effective recommendation systems.

Specifically, this book explores the Apriori Algorithm, one of the fundamental algorithms in data mining, and the basic principles underpinning it regarding association rule mining. Get to know its intricacies, learn to spot traces and correlations in transaction-level data, and uncover potential insights that help drive action. The case study for the Farmer's Market provides a real-life use of the algorithm and shows its power and flexibility.

Whether you are a data science professional or a student advancing on your journey into recommendation systems, this book will provide you with a strong foundation, practical knowledge, and tools for implementation. Use this to guide you through the complicated terrain of association rule discovery and its potential to revolutionise your processes.

Partha Majumdar

Table of Contents

PREFACE ..... **A**

CHAPTER 1: INTRODUCTION TO MARKET BASKET ANALYSIS **1**

1.1 Key Advantages of Market Basket Analysis 3
1.2 Challenges in Market Basket Analysis 3
1.3 Ethical Challenges 4
1.4 Conclusion 5

CHAPTER 2: ASSOCIATION RULE MINING **7**

2.1 Market Basket Transactions 7
2.2 Itemset 8
2.2.1 Types of Itemsets 9
2.2.2 Importance of Itemsets in Association Rule Mining 9
2.2.3 Insights from Itemsets 10
2.3 Support Count (Σ) 10
2.3.1 Importance of Support Count 10
2.3.2 Examples of Support Count 11
2.3.3 Applications of Support Count 11
2.3.4 Insights from Support Count 12
2.4 Support (s) 13
2.4.1 Examples of Support 13
2.4.2 Importance of Support 13
2.4.3 Applications of Support 14
2.4.4 Insights from Support 15
2.5 Frequent Itemset 15
2.5.1 Identifying Frequent Itemsets 16
2.5.2 Types of Frequent Itemsets 16
2.5.3 Importance of Frequent Itemsets 17
2.5.4 Examples of Frequent Itemsets 17
2.5.5 Applications of Frequent Itemsets 17
2.5.6 Insights from Frequent Itemsets 18
2.6 Association Rules 18
2.6.1 Co-Occurrences vs. Causality 19
2.6.2 Interpreting Association Rules 20
2.7 Confidence (c) 20
2.7.1 Examples of Confidence 21
2.7.2 Importance of Confidence 21
2.7.3 Applications of Confidence 22
2.7.4 How Support and Confidence Work Together 22
2.7.5 Insights from Support and Confidence 23
2.8 Association Rule Mining Task 24
2.8.1 Steps in the Association Rule Mining Task 24
2.8.2 Criteria for Rule Selection 25
2.8.3 Practical Example of the Association Rule Mining Task 26

2.8.4 Insights from the Association Rule Mining Task 27
2.8.5 How to Mine Association Rules 27
2.8.6 Understanding Graphically 29
2.9 CONCLUSION 30

CHAPTER 3: APRIORI ALGORITHM 31

3.1 APRIORI PRINCIPLE 31
3.2 ALGORITHM FOR FINDING FREQUENT ITEMSETS 33
3.3 CONCLUSION 34

CHAPTER 4: GENERATING ASSOCIATION RULES 36

4.1 ASSOCIATION RULE GENERATION 36
4.2 PSEUDO-CODE FOR GENERATING ASSOCIATION RULES 38
4.3 FUNCTION TO GENERATE ASSOCIATION RULES 39
4.4 CONCLUSION 40

CHAPTER 5: SUGGESTING ITEMS FOR FARMER'S MARKET 42

5.1 ABOUT THE DATA 42
5.1.1 Imputing Missing Values 44
5.2 MODELLING FOR CREATING A RECOMMENDATION SYSTEM 45
5.3 BUILDING A CLIENT APPLICATION FOR RECOMMENDING PRODUCTS FOR FARMER MARKETS . 53
5.4 CONCLUSION 58

ANNEXURES 60

LIST OF CODE SNIPPETS 60

ABOUT THE AUTHOR I

BOOKS BY THE AUTHOR II

Chapter 1: Introduction to Market Basket Analysis

Market Basket Analysis (MBA) is one of the most fascinating and useful methods used in data analytics. At its core, it predicts sales behaviour based on the items customers buy together. Think of how you walk into a grocery store and grab a loaf of bread. You could also pick up a packet of milk or a pack of butter as you walk the aisles. You are, on some level, without even knowing it, etching metrics and trends in your shopping that reflect business and relationship with the market. Market Basket Analysis refers to finding and analysing such patterns to deliver actionable insights.

Market Basket means items purchased by a customer in the store, presumably coming home in an actual shopping basket. But MBA goes far beyond the traditional shopping basket and is used in many industries, from e-commerce and banking to healthcare and digital marketing. For example, in an online store, MBA helps analyse the virtual cart in which customers add products, and similarly, in a telecom package, MBA helps you find a list of services used. This knowledge enables firms to predict customers' requirements, suggest pertinent products, and boost sales.

Understanding customer preferences in a world where the consumer is the king and everything is competitive is essential. Market Basket Analysis enables organisations to answer key questions: What products do people purchase together? What products go well together? Which items would sell at a higher number if promoted or bundled? Answered well, these questions can revolutionise business. A retailer might, for example, find that customers who purchase pasta are also likely to buy tomato sauce and cheese. With this information, they can create compelling product bundles, recommend complementary items online, or stock these products together in brick-and-mortar stores to trigger impulse buys.

Market Basket Analysis is not limited to the retail domain. In e-commerce, it enables recommendation engines, the systems that recommend other items based on what you've already added to your cart to buy or have bought in the past. In health care, it examines patterns in medical treatment, like which medications are often prescribed together for similar conditions. In manufacturing, a predictive maintenance strategy that utilises MBA helps identify the most common functional fail co-occurrences in a machine's components. It also assists web analytics, allowing website designers to analyse users' paths through a site so they may place content and navigate through a web page optimally.

Market Basket Analysis is exciting because it can identify surprising insights. This is best exemplified by the infamous finding that customers who purchased diapers also tended to buy beer on the same shopping trip. This insight (which at first

glance seems unrelated) resulted from MBA and caused stores to place those items next to each other, greatly increasing both sales. Such findings demonstrate the ability of data to reveal associations that human intuition can fail to capture.

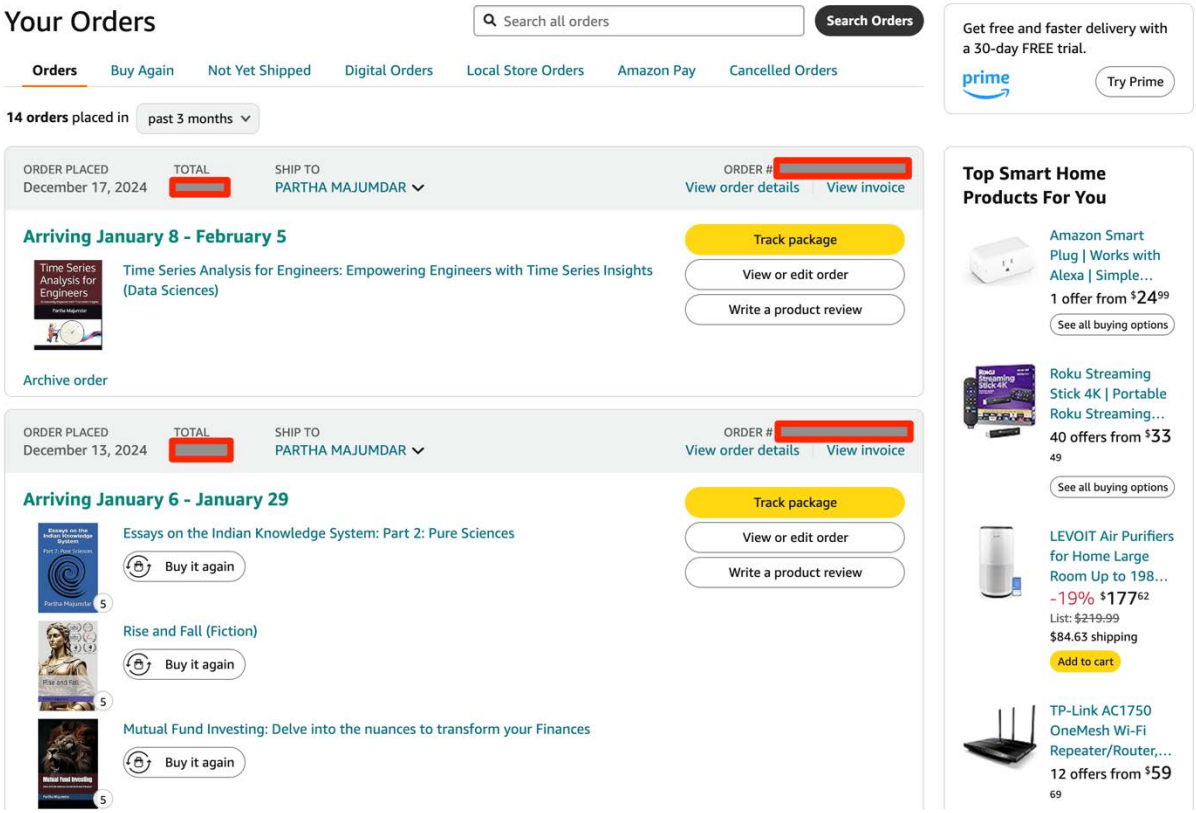

Figure 1: A view of Amazon.com web page showing recommended products.

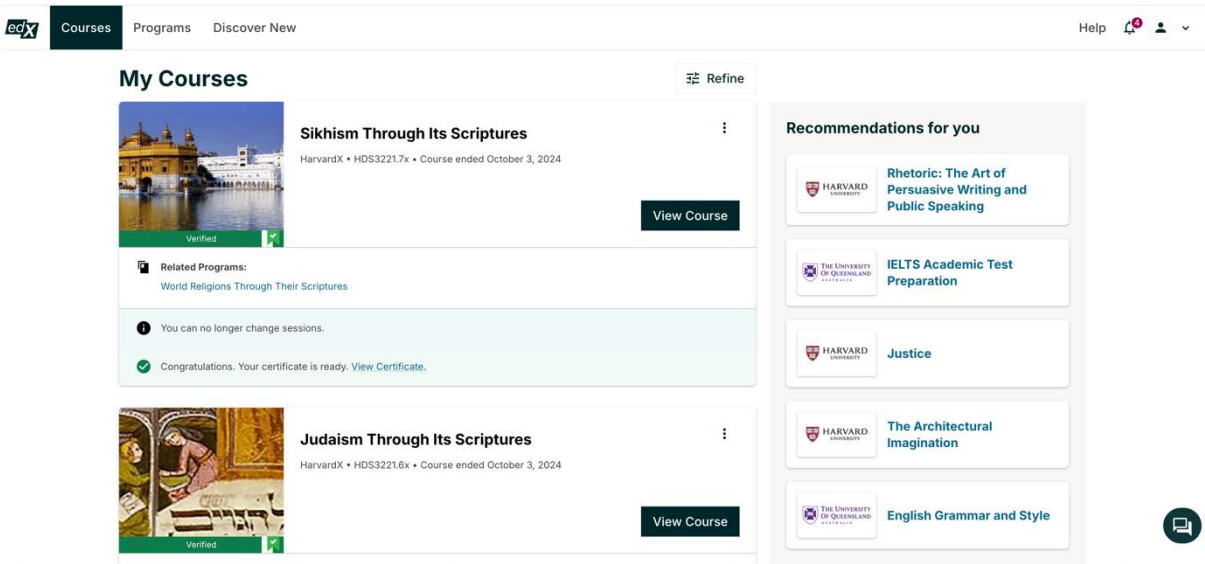

Figure 2: A view of the edx.org web page showing personalised recommended courses.

1.1Key Advantages of Market Basket Analysis

Allowing similar products to be stocked next to each other enables Market Basket Analysis to plan the business's stock effectively. It also improves customer experience by offering personalised product recommendations. It drives marketing campaigns that sync up with customer tastes, such as discounts on packaged goods. It also helps businesses find cross-sell opportunities to convince customers to acquire complementary products to what they already have.

Market Basket Analysis opens up interesting avenues for engineering. Many engineers are instrumental in making MBA a reality. They write algorithms that sift through massive amounts of transactional data, develop efficient storage mechanisms, and design dashboards that give real-time insights. Outside this retail application, engineers apply MBA principles in areas like the Internet of Things (IoT), where data from sensors recognises trends and patterns in machine performance, and cybersecurity, where statistical analysis of network activity can help identify possible threats.

Market Basket Analysis can be applied domain independently hence, it is very popular. In telecommunications, it assists companies in tailoring service bundles that match customer usage patterns. For instance, regular video streaming customers may prefer a plan that bundles high-speed data with streaming platform subscriptions. In banking, it can help identify, for example, which financial products, such as savings accounts and credit cards, are often availed together, enabling banks to create tailored packages. Similarly, the different course enrolments can highlight patterns allowing institutions to build better curriculum structures.

1.2Challenges in Market Basket Analysis

Market Basket Analysis (MBA) has its share of pros and cons despite the benefits it provides. Below is an analysis of these challenges by looking at the nature of transactional data and the complexities of extracting meaningful insights from it.

One major challenge is dealing with the overabundance of data generated in this digital age. Businesses, especially in retail and e-commerce, create huge amounts of transactional data with millions of customers. For instance, global online marketplaces log billions of transactions annually. Stable storage techniques and high-performance computational frameworks are needed to store and analyse these enormous datasets. Traditional tools are often inadequate for dealing with the scale, prompting companies to turn to cutting-edge technologies such as distributed systems and parallel computing.

Another major challenge is finding the signal among the noise. The incidence that two items are often bought together does not mean they are related meaningfully.

For instance, an office customer might buy bread and a novel in one shopping trip in one store, but that may be completely coincidental. Identifying incidental co-purposes versus bona fide relationships is a more sophisticated exercise involving both statistical rigour and qualitative review. If not handled carefully, businesses might use their efforts for false positive patterns, as patterns may not have real-time value.

The sparsity of transactional data adds another level of complexity. Most datasets show that customers only purchase a tiny portion when given a choice of all available items. Such sparsity makes it hard to detect patterns with strong statistical significance, especially in large inventories. In a supermarket containing thousands of unique products, for example, reliable relationships between infrequently purchased items can only be obtained with great analytical sophistication.

The maze of threshold sensitivity in MBA is yet another hurdle. Setting parameters that measure the frequency and strength of associations makes a big difference in the analysis. If these thresholds are set too high, potentially actionable insights will be missed; if they are too low, trivial or irrelevant patterns will swamp the analysis. Finding that balance is often a matter of trial and error, and success depends on the particular goals of the study.

Finally, there is the matter of interpretation of the patterns found, which also has challenges. Although MBA also handles the discovery of these relationships between items, interpreting these relationships does not happen in a vacuum as it would require domain knowledge, especially in the case of two items. In retail, say, an association such as "customers who buy coffee also frequently buy sugar" is intuitive and readily actionable. But in some industries like healthcare or manufacturing, the relationships could be a function of complex interactions between features, like the co-occurrence of certain drugs or the fault patterns of machine parts. Cooperation with domain experts to validate and use this knowledge is essential.

Market Basket Analysis is a great tool for finding hidden patterns between items bought together. However, it comes with challenges that must be addressed to achieve better results. These include effectively scaling the volume of data, extracting signals from noise, dealing with sparse datasets, tuning thresholds correctly, and applying appropriate domain expertise to interpret the results. Overcoming these challenges allows companies to leverage the power of MBA to make better decisions and accomplish their goals more effectively.

1.3Ethical Challenges

Ethical data use is another key aspect of Market Basket Analysis (MBA). Since MBA relies heavily on transactional data to find associations and patterns, organisations

must work to combat any potential problems surrounding customer privacy. Whether for product sales, lead generation, or other purposes, customers trust companies with their data without understanding how it will be used. This creates an obligation for businesses to store and process data fairly, transparently, and securely.

Data ethics alleviate privacy issues and reinforce and preserve customer trust. If customers have a reason to trust that their data will be treated fairly and safely, they are much more likely to interact positively with a brand. Letting customers know what data is being collected, why, and how and giving them the option to decide how their data is used goes a long way in instilling confidence and loyalty.

Moreover, businesses, regardless of size or sector, must comply with data protection mandates like the General Data Protection Regulation (GDPR) in Europe or the California Privacy Rights Act (CPRA) in the United States. Such regulations require companies to collect, store, and process data in a way that respects customer rights and establishes accountability. Following these guidelines helps organisations avoid legal issues and build ethical, customer-centric organisations.

Ethical data practices are the cornerstone of Market Basket Analysis (MBA). While MBA seeks to make businesses profitable, its use should be tightly controlled to ensure insights are not employed unethically, manipulatively, or discriminatorily. Targeted promotions are great if they don't abuse customer weakness or exclude segments.

Market Basket Analysis can be an effective approach that, if done responsibly, can provide valuable data, help drive revenue, and build trust. Through privacy, transparency, and regulatory compliance, businesses can use MBA to gain insight and growth and earn the trust and respect of the customers from whom the data is to be taken.

1.4Conclusion

Market Basket Analysis is both simple and powerful. It combines relations across various items to derive insights from raw data to support value decisions. Whether a neighbourhood business is rearranging its aisles to encourage more impulse buys or a global e-commerce retailer is tuning its recommendation algorithm, it’s all thanks to Market Basket Analysis and the simple insights it provides into the nature of customer behaviour and how best to capitalise on this in the business world.

As you dive deeper into this book, you will discover techniques, tools, and applications of ‘Market Basket Analysis.’ This chapter is contextual, showing you Market Basket Analysis in its true sense and the impact it creates. The upcoming chapters will show you how to apply and leverage this powerful technique, which

gives engineers and analysts greater happiness in gaining value from transactional data.

Chapter 2: Association Rule Mining

Market Basket Analysis's core is a data mining technique called Association Rule Mining. This technique helps businesses identify patterns and relationships in transactional data. It also serves as a basis for gathering important information, such as which items are bought together or how specific customer actions affect others. Association rules lend a formal methodology for assessing massive datasets, taking information from its raw state and converting it into useful information to drive change.

Association Rule Mining aims to discover interesting the "if-then" relationships between the variables in the data. For instance, in a retail context, it could be a rule like, "If a customer purchases bread, then they are likely to buy butter." It is intuitive — and hugely beneficial for businesses looking to improve customer experience, boost sales and streamline operations. Recognising and predicting buying patterns enables organisations to develop targeted marketing strategies, make appropriate product recommendations, and create personalised shopping experiences.

In this chapter, we will study the underlying principles of Association Rule Mining, discussing what it is and how it's used in Market Basket Analysis. We will also discuss how to read and use association rules, applying the theory with real-world cases and practical insights. This chapter lays the groundwork for understanding how rules are generated through important metrics such as support, confidence, and lift and their contributions to mining meaningful patterns.

To transcend industries, we need engineers to harness the power of Association Rule Mining and its concepts. From bettering recommendation systems to improving inventory management or even enhancing customer retention strategies, applications of Association Rule Mining are boundless. This chapter will allow you to discover these possibilities and get you to Market Basket Analysis mastery.

2.1 Market Basket Transactions

Market Basket Transactions are the driving force behind Market Basket Analysis and Association Rule Mining. As per the definition, this relates to the collection of items a customer has purchased together in a single shopping trip or transaction. Each transaction is simply a record of the products that were bought together, and this data provides the raw material for identifying patterns and associations. *As with any data analytics, we need past data, and the Market Basket Transactions are the past data in this case.*

A Market Basket Transaction typically consists of:

- **Transaction ID**: A unique identifier for each transaction.
- **Items Purchased**: A list of products or services included in that transaction.

Analysing these transactions helps uncover relationships, such as items frequently bought together, which can inform strategies for cross-selling, product placement, and personalised recommendations.

The table below illustrates a sample dataset of Market Basket Transactions:

Transaction ID	Items Purchased
1	Bread, Butter, Milk
2	Bread, Eggs
3	Milk, Butter, Cheese
4	Bread, Butter, Jam
5	Cheese, Butter, Bread
6	Milk, Eggs, Bread
7	Bread, Butter, Cheese, Jam
8	Eggs, Milk, Butter

Table 1: Example of Market Basket Transactions.

Each row in the table represents a single transaction. For instance:

- **Transaction 1**: A customer purchased Bread, Butter, and Milk together.
- **Transaction 5**: Bread, Cheese, and Butter were purchased in a single transaction.

From this dataset, patterns and associations can be derived. For example:

- Bread and Butter appear together frequently, indicating a strong association.
- Milk and Eggs are common co-purchases, potentially signifying complementary use.

This simple example sets the stage for understanding how transactional data is processed in Association Rule Mining.

2.2Itemset

The **itemset** is a basic concept in Market Basket Analysis and Association Rule Mining. It refers to a set of one or more products that appear together in a transaction. Itemsets can include single products or combinations bought during one shopping trip. They are the basic units for determining patterns, generating association rules, and discovering correlations from transactional data.

2.2.1 Types of Itemsets

Let us understand the different types of itemsets.

1. Single Itemsets (1-itemsets)

These are individual items, such as Bread, Butter, or Milk, that appear in transactions.

Example: {Bread}, {Milk}, {Butter}

2. Multi-Itemsets (k-itemsets)

These refer to combinations of unique items that occur together. A 2-itemset consists of two unique items, a 3-itemset consists of three, and so on.

Example: {Bread, Butter}, {Milk, Eggs}, {Bread, Milk, Butter}

Examples of Itemsets

We can identify several itemsets from the transactional data in Table 1, as given in Table 2.

Transaction ID	Itemsets
1	{Bread}, {Butter}, {Milk}, {Bread, Butter}, {Bread, Milk}, {Butter, Milk}, {Bread, Butter, Milk}
2	{Bread}, {Eggs}, {Bread, Eggs}
3	{Milk}, {Butter}, {Cheese}, {Milk, Butter}, {Milk, Cheese}, {Butter, Cheese}, {Milk, Butter, Cheese}
4	{Bread}, {Butter}, {Jam}, {Bread, Butter}, {Bread, Jam}, {Butter, Jam}, {Bread, Butter, Jam}
5	{Cheese}, {Butter}, {Bread}, {Cheese, Butter}, {Cheese, Bread}, {Butter, Bread}, {Cheese, Butter, Bread}

Table 2: Itemsets identified from the first five Transactions in Table 1.

2.2.2 Importance of Itemsets in Association Rule Mining

Itemsets have a significant importance as mentioned below.

- **Pattern Recognition**: Businesses can identify combinations of items that tend to be purchased together by analysing frequently occurring itemsets across transactions.

- **Generating Association Rules**: "if-then" relations are built using itemsets. For instance, the frequent itemset {Bread, Butter} can lead to the association rule: “If Bread, then Butter.”
- **Evaluating Metrics**: The itemsets are used to compute metrics such as support, confidence, and lift to measure the strength and relevance of associations.

2.2.3 Insights from Itemsets

We can derive the following as a sample from the examples of itemsets in Table 2.

- The itemset {Bread, Butter} occurs in several transactions, showing a strong association between these products.
- Larger itemsets like {Bread, Butter, Milk} provide valuable insights into bundled purchasing behaviour and highlight opportunities for effective cross-selling strategies.

Therefore, the itemset is the building block of meaningful insights. All organisations must leverage transactional data and derive relationships that inform better decision-making.

2.3Support Count (σ)

Support Count is a primary measure in Association Rule Mining that counts how frequently an itemset occurs in a transaction dataset. It is the initial step in assigning a score for the importance of itemsets, which is then used for advanced metrics like support, confidence, lift, etc.

The Support Count of an itemset is the number of transactions containing the itemset. For instance, if the itemset {Bread, Butter} is present in 4 out of 8 transactions, then Support Count of {Bread, Butter}=4. This metric shows how popular or frequent the itemset is within all transactions.

2.3.1 Importance of Support Count

Support Count is one of the most important tools at our disposal because it allows itemsets that appear infrequently to be disregarded. Businesses can limit their analysis to itemsets with significant relevance or impact on them by establishing a threshold for Support Count (Support). For instance:

- An itemset with a low Support Count may not affect the decision-making.
- High Support Counts itemsets can give actionable insights and serve over strong association rules.

Support Count is also instrumental in determining frequent itemsets (those that satisfy a minimum threshold). This provides the basis for generating association rules.

2.3.2 Examples of Support Count

Using the transactions from **Table 1**, we can calculate the Support Count for various itemsets. To determine the support count, we browse through all the transactions and count the number of times the itemset occurs in all the transactions. The order in which the items in the itemset are placed in a transaction does not matter.

Itemset	Support Count
{Bread}	6
{Butter}	5
{Milk}	4
{Eggs}	3
{Bread, Butter}	4
{Milk, Butter}	3
{Bread, Butter, Milk}	2

Table 3: Examples of Support Count.

Explanation

- {Bread} appears in 6 out of 8 transactions, so its Support Count is 6.
- {Bread, Butter} appears together in 4 transactions, giving it a Support Count of 4.
- {Bread, Butter, Milk} appears together in 2 transactions, making its Support Count 2.

2.3.3 Applications of Support Count

Support Count can be used in the below ways.

Identifying Frequent Itemsets

Support Count determines which are frequent enough for further analysis.

For example, an itemset such as {Bread, Butter} with a high Support Count indicates a strong association worth digging into.

Filtering Irrelevant Data

Establishing a minimum support level (e.g., itemsets must be present in a minimum of 3 transactions) can trim low-frequency itemsets, lessening the complexity of analysis.

Improving Decision-Making

High Support Count itemsets help businesses analyse where customers show the greatest interest. This allows companies to focus on building relevant marketing campaigns, product placements, or bundling strategies.

2.3.4 Insights from Support Count

The support count is not only a basic metric for counting the occurrences of itemsets; it also provides ideas about user behaviour, inventory management, and marketing strategies. Analysing itemsets with high Support Counts enables us to discover patterns that empower us to make smart choices. Support Count is critical in the following areas.

- **Customer Behaviour**: A high Support Count would indicate items frequently purchased together like {Bread, Butter}. This means that both products can work together in a complementary fashion. In the context of businesses, a primary reason for identifying such patterns is that they allow a company to understand its customers better and promote these complementary items together. A grocery store may keep Bread and Butter together or offer a discount if one buys them.
- **Inventory Planning**: Support Count can help businesses optimise stock levels for items often purchased together. For example, Support Count of itemset {Milk, Butter} = 3 implies that milk and butter appear together in 3 transactions. By enabling an understanding of this association, businesses can make better decisions about stocking items in their inventory appropriately to meet consumer demands, which helps to prevent stockouts or overstock.
- **Cross-Selling Opportunities**: Sets with moderate Support Counts like {Bread, Butter, Milk} show cross-selling opportunities. Such itemsets indicate that customers will buy these items together, which can be bundled together for promotional or combo offers. This leaves ample room for sellers to bundle services and maximise Average Transaction Value (ATV) while providing a more pleasant experience for those searching for interconnected products.

Support Count is a numerical value indicating how popular the itemsets are, constituting the basis for producing useful association rules. These insights can enable businesses to identify opportunities for improvement, optimise operations, and develop targeted marketing strategies.

2.4Support (s)

Support is an essential metric in Association Rule Mining that indicates how often an itemset appears in a dataset of transactions. Support is different from Support Count in that Support Count is an absolute value. In contrast, Support is a relative measure that is a percent/proportion of the total number of transactions. Support is calculated as follows.

Support (s) = $\frac{Support\ Count\ (\sigma)}{Total\ Number\ of\ Transactions}$

Support is important because it identifies itemsets in the transactions enough times to be meaningful for analysis. The high-support itemsets reflect the strong patterns of co-occurrence between these items, while the low-support itemsets are often ignored because they yield less significant associations.

Support of an itemset is a number between 0 and 1. A Support of 0 indicates that the itemset does not appear in any transactions. Whereas a Support of 1 indicates that the itemset appears in all the transactions in the dataset.

2.4.1 Examples of Support

Using the transactional data from **Table 1**, we can calculate the Support for various itemsets:

Itemset	Support Count	Total Transactions	Support (s)
{Bread}	6	8	6/8 = 75%
{Butter}	5	8	5/8 = 62.5%
{Milk}	4	8	4/8 = 50%
{Bread, Butter}	4	8	4/8 = 50%
{Milk, Butter}	3	8	3/8 = 37.5%
{Bread, Butter, Milk}	2	8	2/8 = 25%

Table 4: Examples of calculating Support for Itemsets.

2.4.2 Importance of Support

Support is an important concept in Association Rule Mining for these reasons.

Filter out infrequent Itemsets.

Analysis usually does not involve Itemsets with low support. For example, the fact that an itemset occurred in only one out of thousands of transactions is not informative.

Setting a minimum Support threshold allows businesses to eliminate noise and focus on itemsets with some meaningful relevance.

Identifying Popular Patterns

Using Support is a general method to derive interest-provoking transactions.

Example: The itemset {Bread, Butter} has a 50% Support, which means that 50% of all transactions contain the items together and is a strong buying trend.

Generating Association Rules

Support is a basic measure for association rules. A rule should satisfy the minimum Support. This helps us filter irrelevant rules for further analysis of the meaningful ones using metrics like confidence and lift.

2.4.3 Applications of Support

In Association Rule Mining, the Support metric is the first measure that describes the frequency of an itemset in a transaction dataset. Understanding trendy itemsets with high/medium Support enables companies to make smart decisions regarding product placement, marketing campaigns, and inventory management. Here are some of the major uses of Support.

- **Product Placement**: Itemsets with high support can make effective decisions on product placement. If certain items tend to be found together in a transaction, businesses can place one item beside another to make it more accessible for customers and increase sales. For example, if the itemset {Bread, Butter} has a Support of 50%, these products are being bought together in 50% of all transactions. Grocers can place Bread and Butter on the same aisle so shoppers can easily grab it on impulse as they walk by.
- **Marketing Campaigns**: Support plays a significant role in creating effective marketing campaigns. Promotional offers can include combining the products with high or medium Support, encouraging potential products to be cross-sold and increasing the value of the transaction. For instance, if {Milk, Butter} has Support 37.5%, these items are frequently purchased together. By understanding this, businesses can run promotions like discounts or combo deals on Milk and Butter, encouraging customers to buy both products and improving customer satisfaction.
- **Optimising Inventory**: In inventory management, the values of Support are also very important. Frequent demand itemsets denote products a customer is willing to buy, which means they should always be in stock.

The itemset {Bread, Butter, Milk} with a Support of 25% implies that, throughout a quarter of all transactions, this itemset goes jointly. This information allows businesses to prioritise these products in inventory planning, avoiding stockouts and effectively meeting customer demands.

Businesses can leverage the insights drawn from Support to optimise operations, enhance customer experiences, and promote sales growth. Support is a great tool for data-driven decision-making, whether strategically placing products in the right locations, offering attractive promotional bundles, or ensuring stock is managed efficiently.

2.4.4 Insights from Support

Analysing the Support of various itemsets helps businesses find meaningful patterns, improve their strategies, and make data-driven decisions. Here are some key insights extracted from Support values.

- **Identifying Staple Products**: The Support of an individual item can demonstrate its ranking on the marketplace and is a consideration factor for customers. The Support of {Bread} is 75%, meaning that Bread is a staple product purchased in most transactions. This indicates that Bread should never be out of stock and is well placed in the store as it drives customer traffic.
- **Identifying Strong Associations**: With a relatively high Support value of two or more items in an itemset, their combined Support indicates a strong association. For example, the Support for {Bread, Butter} is 50%. This means in 50% of the transactions, Bread and Butter were purchased together. That makes them a prime candidate for cross-selling opportunities or bundled promotions. You can, for instance, sell discounted Bread and Butter when they are bought together.
- **Using Moderately Frequent Combinations**: Conversely, itemsets with moderate Support can reveal potential targeted marketing or placement scenarios. {Bread, Butter, Milk} with 25% Support means this combination is bought in 1 in 4 transactions. It's not as common as staple itemsets, but still an important pattern that businesses can use to their advantage by offering these items as a breakfast set or just putting them nearby in the store for an impulse purchase.

2.5Frequent Itemset

In the context of Association Rule Mining, a Frequent Itemset is a set of items that appear together in a transaction dataset in a way that meets a user-specified minimum Support threshold. These itemsets generate the association rules by

discovering correlations between items in the given dataset. Businesses can use the concept of Frequent Itemsets to reduce the amount of unnecessary information while detecting valuable sets of items that are frequently purchased in transactions.

Frequent Itemsets are important tools in Market Basket Analysis. They signify the items or groups that tend to be purchased together. Based on the data analysis, this can provide businesses with insights about product placement, inventory management, and promotions.

2.5.1 Identifying Frequent Itemsets

Frequent Itemsets are derived by calculating each itemset's Support in the transaction dataset. Only those itemsets whose Support meets or exceeds the minimum Support threshold are considered frequent.

For example, using **Table 1**, if the minimum Support threshold is 50%, the following itemsets qualify as Frequent Itemsets.

- {Bread} with Support of 75%
- {Butter} with Support of 62.5%
- {Bread, Butter} with Support of 50%

Itemsets with lower Support, such as {Bread, Butter, Milk} with 25% Support, would not qualify as Frequent Itemsets under this threshold.

2.5.2 Types of Frequent Itemsets

We can classify Frequent Itemsets as of two types.

Single Frequent Itemsets

- These consist of individual items that meet the minimum Support threshold.
- Example: {Bread} and {Butter} are single Frequent Itemsets with Support above 50%.

Frequent Multi-Itemsets

- These are combinations of items that co-occur frequently enough to meet the threshold.
- Example: {Bread, Butter} is a Frequent Multi-Itemset with Support 50%.

2.5.3 Importance of Frequent Itemsets

Frequent Itemsets are key in Association Rule Mining for the following reasons.

- **The basis for association rules**: Frequent Itemsets are used in generating the association rules that assist businesses in analysing the relationships between items. For example, the Frequent Itemset {Bread, Butter} will lead to the Rule: "If Bread, then Butter."
- **Pattern Recognition**: Frequent Itemsets are combinations of items customers frequently buy together. They help businesses identify patterns and customer preferences.
- **Efficient Filtering**: By focusing on Frequent Itemsets, businesses can eliminate the noise caused by infrequent itemsets, making the analysis more straightforward and ensuring that only significant patterns are analysed.

2.5.4 Examples of Frequent Itemsets

The following Frequent Itemsets are identified using **Table 1** and a minimum Support threshold of 50%.

Itemset	Support Count	Total Transactions	Support
{Bread}	6	8	75%
{Butter}	5	8	62.5%
{Bread, Butter}	4	8	50%

Table 5: Frequent Itemsets based on the Market Basket Transactions in Table 1.

These itemsets qualify as frequent because their Support meets or exceeds the 50% threshold.

2.5.5 Applications of Frequent Itemsets

Frequent Itemsets help to understand the actionable aspects which can be implemented in a business.

- **Product Placement**: Companies can effectively group frequently purchased products by analysing frequent item sets. For example, placing Bread and Butter appropriately encourages customers to buy.
- **Promotional Campaigns**: All Frequent Itemsets, such as {Milk, Butter}, can be used to create combo offers or discounts, which can drive customers to buy these together.
- **Inventory Management**: Frequent Itemsets identify products in regular demand, so we ensure these items are not out of stock in inventory planning.

2.5.6 Insights from Frequent Itemsets

Finding Frequent Itemsets helps businesses generate meaningful insights.

- **Understanding Customer Behaviour**: From the above Frequent Itemset, we infer that Bread is a staple bought in 75% of the transactions. Thus, Bread is an important item to stock.
- **Finding Strong Relations**: {Bread, Butter} correlates well since it appears in 50% of transactions. This can mean an opportunity for cross-selling or bundle promotions.
- **Targeted Marketing Opportunities**: Even moderately frequent itemsets, such as {Bread, Butter, Milk}, can be used for targeted marketing strategies (e.g., breakfast bundles).

An essential aspect of Association Rule Mining is the design of frequent itemsets, allowing businesses to concentrate on critical trends in transactional information. An organisation can focus their operations, building up its customer satisfaction and bringing the business's overall performance to the best.

2.6Association Rules

Association Rules are "if-then" statements that specify the relationship between items in the transactional data in the Market Basket Analysis. These rules are to find and report where the two items tend to co-occur, such as "If you bought Bread, then you're likely to buy Butter". This is all done through analysing the co-occurrences of items in multiple transactions.

For example, we get the following from the transactional data in Table 1.

- **Rule 1**: *If Bread, then Butter.*
- **Rule 2**: *If Milk, then Eggs.*
- **Rule 3**: *If Milk & Bread, then Eggs.*

We write this as follows.

- **Rule 1**: *{Bread} → {Butter}*
- **Rule 2**: *{Milk} → {Eggs}*
- **Rule 3**: *{Milk, Bread} → {Eggs}*

On the left-hand side (like {Bread}, {Milk}, and {Milk, Bread}) are the items the customer has purchased. On the right-hand side (like {Butter}, {Eggs}, {Eggs}) are the items that can be recommended to the customer.

These rules indicate that when customers buy certain items, they'll likely buy others in the same purchase. They are extremely useful for businesses to improve product recommendation systems, optimise inventory positioning, and devise better cross-sell strategies.

2.6.1 Co-Occurrences vs. Causality

One important part of Association Rule Mining is that it finds co-occurrences rather than causation. This is important to understand what information you would get from the Market Basket Analysis.

2.6.1.1 What Are Co-Occurrences?

Co-occurrence refers to the items that appear together most frequently in a transaction. In our example:

- Bread and Butter are frequently seen together in the same transaction. So, they form an association rule.
- The reason for this pattern may be that they are complementary products (e.g. they are used together frequently for breakfast)

2.6.1.2 What Is Causality?

Causality means that one event directly causes another. For example, a causal relationship exists if buying Bread leads the customer to purchase Butter. However, Association Rule Mining does not show this direct cause-effect relationship. It only tells us that Bread and Butter are frequently bought together.

One requisite of causality is temporal information, which we do not track in the Market Basket Transactions.

2.6.1.3 Why Co-Occurrences Are Not Causal

Co-occurrences are not causal for the following reasons.

- **Correlation, Not Causation**: Association rules measure correlation - how frequently two items co-occur. If successful, they bring no explanation for the relationship. For instance, Bread and Butter might co-occur because they are often part of a meal but do not necessarily directly "cause" the purchase of each other.
- **Impact of External Influencers**: Third-party interactions such as product placements or promotions could trigger co-occurrences. For example, if Bread and Butter are next to each other in a store, a customer is likely to pick up both Bread and Butter, creating a pattern of co-occurrence by which there is no direct causation.
- **Customer Behaviour**: Customer preferences, habits or recipes might drive such patterns. For example, a customer who purchased Bread may also want Butter to make for toast, but that is a need driven by the customer's intent and not the purchase of Bread itself.

- **Incidental Relationships**: In some cases, the co-occurrences could be coincidences. For instance, if Bread and Jam are found together in a few transactions, it might not be because of any interesting relation but because of a small subset of customers.

2.6.2 Interpreting Association Rules

While association rules can be extremely helpful in discovering new opportunities, they must be interpreted carefully.

- They illuminate what occurs together in transactions, not why.
- Businesses must develop these insights with domain knowledge to uncover customer motivations and frame effective strategies.

2.7Confidence (c)

Confidence is a basic metric of Association Rule Mining, which describes how often the item in question (or set of items) is purchased, given that another item was also purchased. It indicates the strength or correlation of an association rule. It is useful for businesses to know the likelihood that the "then" part of the association rule will happen if the "if" part is true.

Confidence is calculated as follows.

Confidence = $\frac{Support\ of\ (Antecedent, Consequent)}{Support\ of\ Antecedent}$

Or, c = $\frac{s(Antecedent, Consequent)}{s(Antecedent)}$

Or, c = $\frac{\sigma(Antecedent, Consequent)}{\sigma(Antecedent)}$

Where:

- Antecedent: The "if" part of the rule.
- Consequent: The "then" part of the rule.

For example, if the rule is **{Bread}** → **{Butter}**, Confidence answers the question: "What is the probability that Butter is purchased when Bread is purchased?"

Confidence is a value between 0 and 1 because the Support for the Antecedent and the Consequent is always less than or equal to the Support for the Antecedent. This is because the joint occurrences of the Antecedent and Consequent cannot exceed the occurrences of the Antecedent.

2.7.1 Examples of Confidence

Using the transactional data from **Table 1**, let's calculate the Confidence for some association rules.

1. **Rule: {Bread} → {Butter}**

   Support of {Bread, Butter} = 50%

   Support of {Bread} = 75%

   Confidence = $\frac{50\%}{75\%}$ = 66.67%

   Interpretation: If Bread is purchased, there is a 66.7% chance that Butter will also be purchased.

2. **Rule: {Milk} → {Butter}**

   Support of {Milk, Butter} = 37.5%

   Support of {Milk} = 50%

   Confidence = $\frac{37.5\%}{50\%}$ = 75%

   Interpretation: If Milk is purchased, there is a 75% chance that Butter will also be purchased.

3. **Rule: {Bread, Butter} → {Milk}**

   Support of {Bread, Butter, Milk} = 25%

   Support of {Bread, Butter} = 50%

   Confidence = $\frac{25\%}{50\%}$ = 50%

   Interpretation: If Bread and Butter are purchased, there is a 50% chance that Milk will also be purchased.

2.7.2 Importance of Confidence

Confidence is the most important property of association rules and is used in analysis.

- **Strength of Relationships**: Confidence measures the strength of the relationship between Antecedents and Consequent. Larger Confidence values suggest more robust patterns.

- **Rule Validation**: Confidence ensures only meaningful and reliable rules are being analysed. Rules with very low Confidence are often ignored because they do not always hold.
- **Actionable Insights**: Using Confidence, prediction rules can guide better scheme decisions. For example, a strong rule such as {Bread} → {Butter} can help place those products together in the supermarkets or assist in determining promotions.

2.7.3 Applications of Confidence

Confidence is a crucial component for many business purposes.

- **Product Recommendations**: Recommendations engines can be driven using high-confidence rules. If the rule {Bread} → {Butter} has a Confidence of 66.7%, then it is likely that selling Butter to customers who buy Bread will be successful.
- **Promotional Strategies**: With confidence, understand what products to bundle or cross-sell. Example – A rule {Milk} → {Butter} with 75% Confidence means if Customer buys Milk, give him a discount on Butter
- **Inventory Planning**: Confidence gives predictions about demand for related goods. An association rule {Bread, Butter} → {Milk} with 50% Confidence suggests that {Milk} can be stocked near product {Bread} and {Butter} in stores, since if users are buying {Bread, Butter} → {Milk}, stores can fulfil consumer needs.

2.7.4 How Support and Confidence Work Together

We use Support and Confidence metrics to measure the quality of association rules. They are closely related.

Support shows how often items in the rule appear in the dataset. A rule with High Support indicates sufficient data to generate and ensure it is not rare.

Confidence looks at the reliability of the rule, i.e., the Consequent's probability when the Antecedent is true.

Examples:

For the rule **{Bread}** → **{Butter}**:

- Support is 50% (meaning that 50% of all transactions result in purchasing Bread and Butter together).
- Confidence = 66.7% (meaning that Butter is purchased 66.7% of the time Bread is purchased).

High Support guarantees a statistically significant rule, while high confidence guarantees reliability. Typically, a good rule has both high Support and Confidence.

2.7.5 Insights from Support and Confidence

Let us discuss some valuable insights from the combination of Support and Confidence.

1. Strong Associations

High Support indicates the rule appears in many transactions. High Confidence indicates the item is likely present when you have the first item. These rules indicate trends that are common and reliable in predicting customer actions.

Example:

- **Rule**: {Milk} → {Butter}
- **Support**: 37.5% (showing Milk and Butter occur together 37.5% of the time in transactions).
- **Confidence**: 75% (If someone buys Milk, they will also buy Butter in 75% of cases).

This rule indicates a definite relationship between Milk and Butter, which can be valuable in deciding product placement or bundled promotions.

2. Targeted Strategies

High Support and Confidence are best, but low Support and high Confidence can still allow you to market products as niche. These rules can suggest certain less frequent combinations are still valid customer needs or behaviours.

Example:

- **Rule**: {Bread, Butter, Milk} → {Jam}
- **Support**: 12.5% (this combination occurred in only 12.5% of the transactions).
- **Confidence**: 80% (wherein, if Bread, Butter, and Milk were bought, Jam carries an 80% chance of being bought as well).

This rule shows sideways play opportunities. They can help drive sales in a specific segment. For example, a business could promote Jam with Bread, Butter, and Milk and offer it at a discount.

3. Business Decision-Making

Combining Support and Confidence allows businesses to balance statistical significance (through Support) and reliability (through Confidence). Such a design has a high probability of yielding good enough quality rules with a pertaining meaning, and actions can be taken based on the derived rules.

For instance, a rule with high Support but low Confidence may not be reliable enough. Conversely, a rule with a low Support score yet a high Confidence score may indicate an underutilised opportunity for a specific customer group.

By analysing those two metrics, businesses can:

- Create successful marketing campaigns.
- Optimise product placement.
- Create group promotions and/or loyalty programs.
- Improve inventory management strategies to fulfil customer demand better.

So, combined Support and Confidence comprehensively measure how well an association rule describes the data. Rules in which both metrics attain relatively high values are considered strong and reliable patterns, while low Support with high Confidence represents a niche opportunity.

Armed with these insights, businesses can:

- Pay attention to the right patterns.
- Optimise their operations.
- Improve customer satisfaction.

By evaluating Association Rules using both Support and Confidence from the bottom up, we can guarantee that they will be statistically significant between the Rule, Common Sense, and impactful decision-making.

2.8Association Rule Mining Task

Association Rule Mining is used to find important relationships in large datasets. The rules that satisfy minimum Support and minimum Confidence rules are statistically significant and reliable. They are an essential part of Market Basket Analysis and have various applications across data-driven sectors.

2.8.1 Steps in the Association Rule Mining Task

The steps in Association Rule Mining are as follows.

1. Identify Frequent Itemsets

Identify all itemsets that fulfil the minimum Support threshold set by the user.

For instance, consider the transaction details provided in Table 1; if the minimum Support threshold is 50%, then we recognise that {Bread}, {Butter}, and {Bread, Butter} are Frequent Itemsets with their respective Support values being 75, 62.5%, and 50%.

2. Generate Association Rules

After discovering Frequent Itemsets, the next step is to generate association rules from them. Each rule is represented as A → B, with the Antecedent (Rules left-hand side items or A) and the Consequent (Rules right-hand side items or B).

For example, The rules *Bread* → *Butter* and *Butter* → *Bread* can be derived from the Frequent Itemset {Bread, Butter}.

3. Evaluate Rules against the Minimum Support and Confidence.

The common approaches for rule evaluation are using Support and Confidence.

Each rule is checked against the minimum Confidence threshold that it is reliable.

The rule *Bread* → *Butter* has a Confidence of 66.7%, which beats a minimum threshold of 60%.

The minimum Support and minimum Confidence are set based on the business context.

2.8.2 Criteria for Rule Selection

The following are the criteria for Association Rule selection.

Minimum Support

Rules must satisfy the minimum Support threshold, ensuring the pattern is statistically significant and not based on rare occurrences.

Example: If the Support for *Bread* → *Butter* is 50%, and the minimum threshold is 40%, this rule is included for further evaluation.

Minimum Confidence

Rules must meet the minimum Confidence threshold to ensure the reliability of the association.

Example: If the rule *Milk* → *Butter* has a Confidence of 75% and the threshold is 60%, the rule qualifies.

2.8.3 Practical Example of the Association Rule Mining Task

Using Table 1, let's identify associations based on the following thresholds:

- Minimum Support: 50%
- Minimum Confidence: 60%

Frequent Itemsets:

- {Bread} with Support 75%
- {Butter} with Support 62.5%
- {Bread, Butter} with Support 50%

Generated Rules

- Rule 1: *Bread* → *Butter*
- Rule 2: *Butter* → *Bread*

Evaluate Rules

Rule 1: *Bread* → *Butter*

- Support: 50%
- Confidence: $\frac{\sigma(Bread,Butter)}{\sigma(Bread)} = \frac{4}{6} = 66.67\%$
- Thus, Valid.

Rule 2: *Butter* → *Bread*

- Support: 50%
- Confidence: $\frac{\sigma(Bread,Butter)}{\sigma(Butter)} = \frac{4}{5} = 80\%$
- Thus, Valid.

Both rules are statistically significant and reliable, meeting the thresholds for Support and Confidence.

2.8.4 Insights from the Association Rule Mining Task

Here are some insights we could get from the above Association Rule Mining Tasks.

- **Strong Associations**: Combining support and confidence yields rules like *Bread* → *Butter* (Support: 50%, Confidence: 66.7%), which showcase strong customer buying patterns. These patterns can help decide product placement, bundled offers, etc.
- **Targeted Opportunities**: Low Support, High Confidence rules like {*Bread, Butter*} → *Jam*, indicating potential niche marketing opportunities.
- **Business Impact**: Support and Confidence can be combined to give businesses a statistically driven set of rules. These rules lead to effective decisions about marketing strategies, inventory management, and improving customer experience.

2.8.5 How to Mine Association Rules

The simplest method to mine Association Rules is to apply the brute force method. In this technique, we do the following.

1. List all the possible Association Rules for a given Market Basket Transaction dataset.
2. Find the Support and Confidence for each Association Rule.
3. Retain the Association Rules that meet the minimum Support and minimum Confidence criteria.

While this process works, it is computationally very expensive. Let us get an idea of this complexity.

Suppose our Market Basket Transactions contain d unique items. Then the number of itemsets we can form is $\sum_{k=0}^{d}\binom{d}{k} = 2^d$. This implies that we could define $\sum_{k=1}^{d-1}[\binom{d}{k} * \sum_{j=1}^{d-k}\binom{d-k}{j}] = 3^d - 2^{d+1} + 1$ association rules.

We could write this as follows.

Number of unique items = d

Number of itemsets = 2^d

Number of possible Association Rules =

$\sum_{k=1}^{d-1}[\binom{d}{k} * \sum_{j=1}^{d-k}\binom{d-k}{j}] = 3^d - 2^{d+1} + 1$

This formula can be understood as follows: If we have d unique items in our Market Basket Transaction dataset and try to make associations of the type X → Y, we

can choose k items in X in ${}^{d}C_{k}$ ways. Then, there can be ${}^{d-k}C_{j}$ ways to select the items in Y. We have to consider all possible values of k and j.

If our Market Basket Transactions contain six unique items, there will be 64 itemsets and $3^{6} - 2^{6+1} + 1 = 602$ possible association rules.

If our Market Basket Transactions contain ten unique items, the number of itemsets will be 1024, and the number of possible association rules would be $3^{10} - 2^{10+1} + 1 = 57{,}002$.

It is easy to see that finding the frequent itemsets by finding the Support and comparing it with the minimum Support is a computationally intensive task in itself.

Notice how the number of possible association rules grows exponentially as the number of unique items in our Market Basket Transactions dataset grows. Let us plot this using a Python program.

```
import numpy as np
import matplotlib.pyplot as plt

# Define the range of the variable 'd'
d = np.arange(0, 21, 1)  # Integers from 0 to 20

# Define the function f(d) for the number of possible association rules
def f(d):
    return 3**d - 2**(d+1) + 1

# Compute the y-values for the defined integer points
y = f(d)

# Interpolate between points for a smooth curve
d_smooth = np.linspace(0, 20, 500)  # Generate 500 points between 0 and 20
y_smooth = f(d_smooth)

# Plot the function
plt.figure(figsize=(10, 6))  # Set the figure size
plt.plot(d_smooth, y_smooth, label=r'$f(d) = 3^d - 2^{d+1} + 1$', color='blue', linewidth=2)  # S
mooth line
plt.scatter(d, y, color='red', label='Integer Points', zorder=5)  # Overlay integer points

# Additional plot settings
plt.title("Number of Possible Association Rules", fontsize=14)
plt.xlabel("Number of Unique Items (d)", fontsize=12)
plt.ylabel("Number of Possible Rules (f(d))", fontsize=12)
plt.grid(True, linestyle='--', alpha=0.6)  # Add grid lines
plt.xticks(ticks=np.arange(0, 21, 1))
plt.legend(fontsize=12)  # Show legend
plt.tight_layout()  # Adjust layout for better appearance

# Show the plot
plt.show()
```

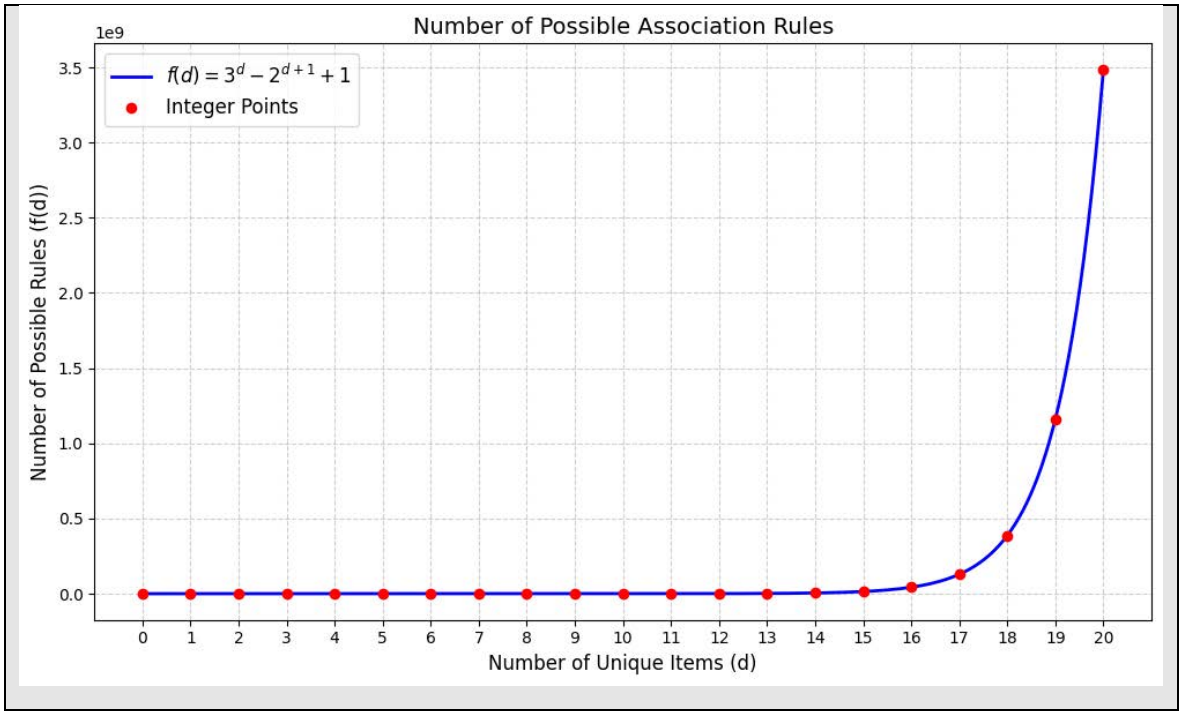

Code 1: NumberOfAssociationRules.ipynb (https://drive.google.com/file/d/1VJyr60qxD7lOe48edePkAK2BRGIO3qTE/view?usp=sharing)

When d = 20, the number of possible association rules = 3,48,46,87,250 (~3.5 billion). This is a huge number. In real life, an online retail store will have millions of products. So, the number of possible association rules will be unmanageable. So, we need smarter ways of dealing with this.

2.8.6 Understanding Graphically

Let's understand this graphically. Our Market Basket Transactions contain five unique items. As shown in the figure below, we can form the different itemsets graphically.

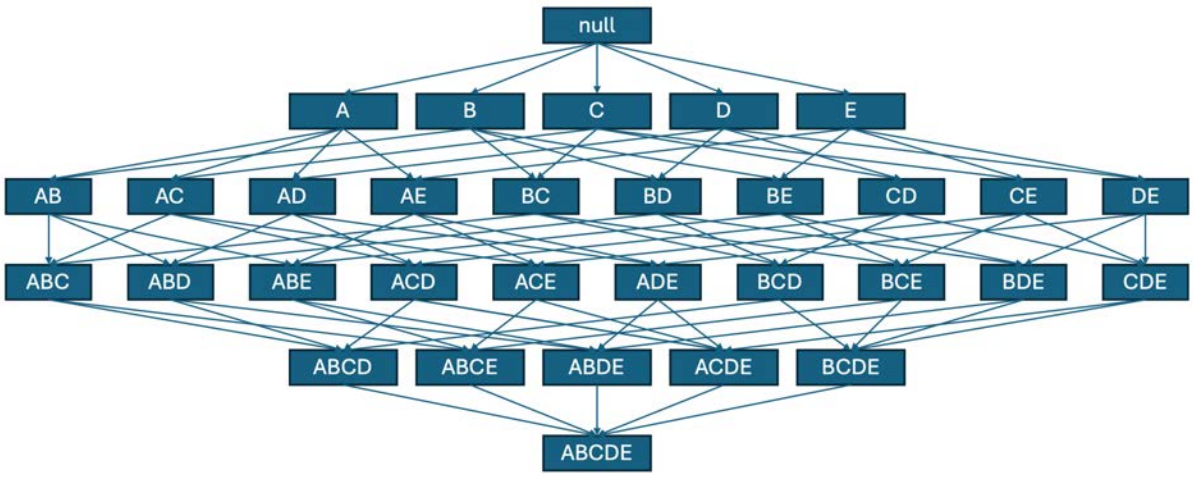

Figure 3: Possible Itemsets from five unique items.

We get as follows.

Number of Itemsets with Zero Items = ${}^5C_0 = \frac{5!}{0!*5!} = 1$

Number of Itemsets with 1 Item = ${}^5C_1 = \frac{5!}{1!*4!} = 5$

Number of Itemsets with 2 Items = ${}^5C_2 = \frac{5!}{2!*3!} = 10$

Number of Itemsets with 3 Items = ${}^5C_3 = \frac{5!}{3!*2!} = 10$

Number of Itemsets with 4 Items = ${}^5C_4 = \frac{5!}{4!*1!} = 5$

Number of Itemsets with 5 Items = ${}^5C_5 = \frac{5!}{5!*0!} = 1$

So, Total number of Itemsets = ${}^5C_0 + {}^5C_1 + {}^5C_2 + {}^5C_3 + {}^5C_4 + {}^5C_5 = 1 + 5 + 10 + 10 + 5 + 1 = 32 = 2^5$

So, the number of Itemsets = $2^d = 2^5 = 32$.

2.9Conclusion

In this chapter, we explored the foundational concepts of Association Rule Mining, which lies at the heart of Market Basket Analysis. We delved into the intricacies of itemsets, support, confidence, and the metrics used to uncover hidden patterns in transactional data. From understanding frequent itemsets to evaluating association rules, the chapter highlighted the computational challenges and the exponential growth of possible rules as the dataset grows.

Association Rule Mining empowers businesses to derive actionable insights, improve product placement, optimise inventory, and design targeted marketing campaigns. However, as the number of possible rules grows rapidly with the dataset's complexity, smarter techniques are essential to manage this scale.

The next chapter will introduce the Apriori Algorithm, a cornerstone in association rule mining that efficiently generates frequent itemsets and overcomes computational limitations. It will help readers understand how to mine association rules effectively, ensuring meaningful and actionable results in large datasets.

Chapter 3: Apriori Algorithm

The Apriori algorithm is one of the fundamental algorithms of data mining, particularly in association rule mining and Market Basket Analysis. It is a systematic approach for finding frequent patterns in large data sets. With these principles, companies and analysers are expected to make useful connections between the items used or bought, creating interpretations that can be utilised in many fields.

We built this algorithm by performing an iterative search for itemsets to meet strong minimum support. It starts by searching for singletons and then expands to larger combinations, each time maintaining computational efficiency by minimising the combinations that did not reach the threshold during earlier iterations. This systematic pruning makes Apriori so efficient and scalable—capable of handling databases with millions of transactions.

The Apriori algorithm is essential for solving the computational problem associated with association rule mining. After you train association rule mining on data, datasets scale massively, so old-school methods don't apply. The algorithm's primary approaches address this shortcoming through a divide-and-conquer method that minimises the effort in producing and validating the itemsets. This technique allows analysts to focus on real trends while gating noise and redundancies.

In this chapter, we will discuss how the Apriori algorithm works with step-by-step technical details with examples.

3.1 Apriori Principle

The Apriori principle is the foundational logic behind the Apriori algorithm. The principle is simple but effective; **if an itemset is infrequent, all of its supersets will be infrequent as well**. This can also be stated as: **If an itemset is frequent, all of its subsets must also be frequent**. This logical inference enables the detection of frequent itemsets that prune the search space, thus leading to a massive improvement in computational power efficiency.

In order to understand this principle, let's take an example of a dataset in supermarket transactions. Let the minimum support threshold be given by α and let the itemset {A, B} be found in fewer transactions than α. Using the Apriori principle, we can see that no superset of {A, B}, like {A, B, C}, {A, B, D, E}, etc., can be frequent as the addition of more items would either reduce or at least keep the original itemset's support same. As a result, we don't have to compute these supersets, which saves computation time and resources.

These rules stem from the **anti-monotone property of support** (i.e., that the support of an itemset cannot exceed the support of any of its subsets). That is, the support of {Milk, Bread, Butter} will always be less than or equal to the support of {Milk, Bread} or of {Milk}. As we make use of this property, the Apriori algorithm eliminates needless calculations for infrequent itemsets and only considers promising candidates in the next iterations.

Mathematically, we can write this as follows, considering the itemsets X and Y.

$\forall X, Y: (X \subseteq Y) \Rightarrow s(X) \geq s(Y)$

This is read as: For all the itemsets X and Y, where X is a subset of Y, the support of X is greater than or equal to the support of Y.

Let us visualise this.

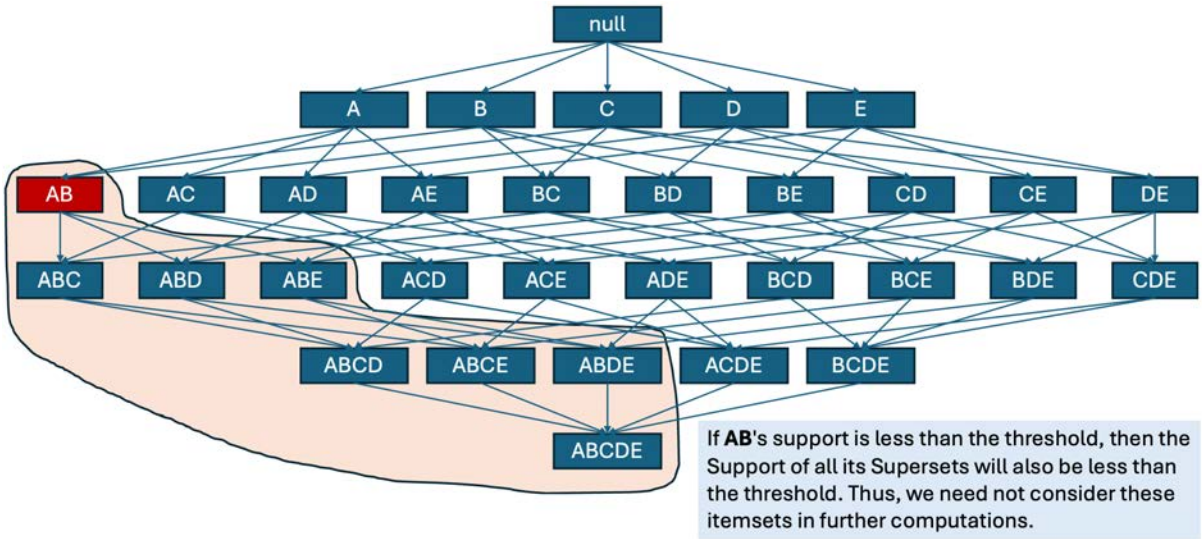

*Figure 4: Illustration of Apriori Principle: If **AB**'s support is less than the threshold, then the Support of all its Supersets (shown in the highlighted portion) will also be less than the threshold. Thus, we need not consider these itemsets in further computation for determining frequent itemsets.*

The above illustration shows that using the Apriori principle, we can drastically reduce the number of itemsets to evaluate to determine the frequent itemsets.

The usefulness of the Apriori Principle is especially visible when examining large datasets. Transaction data in real-world applications such as Market Basket Analysis usually contains thousands of items and millions of records. It would be computationally infeasible to compute frequent patterns if we don't trim the search space by using the Apriori Principle. The inclusion of this principle means that rather than looking at all possible combinations of items (which are potentially in the billions) for being frequent, it filters out virtually all combinations that are unlikely to be considered and narrows it down to just those sets that have an actual chance of being frequent in the dataset, thus simplifying the problem significantly.

The Apriori Principle is the foundation of the Apriori algorithm, which helps find candidate itemsets for evaluation. The result is the simplicity of the steps involved, which translates into better computational performance and scalability of the algorithm with size and complexity. This establishes this algorithm as a major one employed in association rule mining and beyond.

3.2Algorithm for Finding Frequent Itemsets

Given below is the pseudo code for implementing the Apriori Algorithm for finding the frequent itemsets.

1. **Initialization**
   - Begin with identifying all individual items in the dataset, termed as 1-itemsets.
   - Compute the support count for each 1-itemset.
   - Retain only those 1-itemsets whose support count meets or exceeds the minimum support threshold. These are the frequent 1-itemsets (F1).
2. **Iterative Candidate Generation**
   - Set k = 1.
   - Generate candidate (k+1)-itemsets (L_{k+1}) by joining frequent k-itemsets (F_k) with one another.
   - Use the Apriori Principle: prune any candidate (k+1)-itemset if it contains subsets of length k that are not frequent.
3. **Support Counting**
   - For each candidate (k+1)-itemset in L_{k+1}, calculate its support by scanning the database and counting the transactions where the itemset occurs.
4. **Candidate Elimination**
   - Discard candidates in L_{k+1} that do not meet the minimum support threshold. The remaining candidates form the frequent (k+1)-itemsets (F_{k+1}).
5. **Termination**
   - Repeat the process (steps 2–4) until no new frequent itemsets are generated, i.e., F_{k+1} is empty.
6. **Output**
   - Compile all frequent itemsets across iterations (F_1, F_2, …, F_k).

Below is the Python implementation of the Apriori Algorithm.

```
from itertools import combinations

def generate_candidates(L, k):
    return set([frozenset(sorted(set(i) | set(j))) for i in L for j in L if len(set(i) |
 set(j)) == k])

def apriori(transactions, min_support):
    items = set(item for transaction in transactions for item in transaction)
    itemset = set(frozenset([item]) for item in items)
    L = [set()]
    k = 1

    while itemset:
        counts = {i: sum([i.issubset(transaction) for transaction in transactions]) for
i in itemset}
        L_k = set(i for i in counts if counts[i] >= min_support)
        L.append(L_k)
        k += 1
        itemset = generate_candidates(L_k, k)
```

```
    return L[1:]

# Dataset from Table 1
transactions = [
    {'Bread', 'Butter', 'Milk'},
    {'Bread', 'Milk'},
    {'Butter', 'Milk'},
    {'Bread', 'Butter'},
    {'Bread', 'Cheese', 'Butter'},
    {'Milk', 'Cheese'},
    {'Cheese'},
    {'Bread', 'Milk', 'Butter'}
]

min_support = 4  # 50% of 8 transactions

frequent_itemsets = apriori(transactions, min_support)

print("Frequent Itemsets:")
for k, itemsets in enumerate(frequent_itemsets, 1):
    print(f"{k}-itemsets:")
    for itemset in itemsets:
        print(f"  {set(itemset)}")

Frequent Itemsets:
1-itemsets:
  {'Butter'}
  {'Bread'}
  {'Milk'}
2-itemsets:
  {'Bread', 'Butter'}
```

Code 2: AprioriAlgorithm.ipynb (https://drive.google.com/file/d/185F4tXmz6Ao7gu7LQAPFQzkeUNGnYRe1/view?usp=sharing)

3.3Conclusion

The Apriori Algorithm is essential in Association Rule Mining and Data Analysis, especially in Big Data. Efficient due to the use of the Apriori Principle, which allows us to discard itemsets that possess low support itemsets early on in the process. This characteristic ensures that the algorithm is scalable and thus enables its usage in real-time settings such as Market Basket Analysis, Healthcare Analytics, etc.

Employing an iterative algorithm that systematically prunes candidate itemsets guarantees both accuracy and computational efficiency. Minimally supported itemsets readily reveal the algorithm's desire to find useful patterns and relationships in the data. The algorithm's pristine relevance in modern data mining approaches is accentuated by its ability to be utilised in varying areas.

This chapter explained how the Apriori algorithm discovered frequent item sets, creating a solid foundation for understanding. This knowledge forms the basis for our subsequent exploration of how to transform these into association rules that provide insight. In the following chapter, we explore how to generate those

associations and turn the frequent item sets reported by the Apriori Algorithm into actionable suggestions for decision-making.

Chapter 4: Generating Association Rules

The generation of association rules is a crucial phase of association rule mining and Market Basket Analysis. Once these frequent itemsets are discovered, the next step is to find the associations between items in these sets. Association rules map this relationship to "if-then" statements that reveal actionable insights into customer behaviour, product affinities, and market dynamics.

4.1 Association Rule Generation

We generate the rules after we have identified the frequent itemsets. Suppose our frequent itemset is L. First, we need to find all non-empty itemsets f such that $f \subset L$ and $f \rightarrow (L-f)$.

If {A, B, C, D} is our frequent itemset, then the candidate rules will be as follows.

{A} → {B, C, D}

{B} → {A, C, D}

{C} → {A, B, D}

{D} → {A, B, C}

{A, B} → {C, D}

{A, C} → {B, D}

{A, D} → {B, C }

{B, C} → {A, D}

{B, D} → {A, C}

{C, D} → {A, B}

{A, B, C} → {D}

{A, B, D} → {C}

{A, C, D} → {B}

{B, C, D} → {A}

In this case, the number of rules we generated is ${}^4C_1 + {}^4C_2 + {}^4C_3 = \frac{4!}{1!*3!} + \frac{4!}{2!*2!} + \frac{4!}{3!*1!} = 4 + 6 + 4 = 14 = 2^4 - 2$.

We can generalise this to mean that the number of candidate rules we can generate from a k-itemset is (2^k – 2). Notice that we discard the trivial rules {} → {A, B, C, D} and {A, B, C, D} → {}.

Let us visualise this.

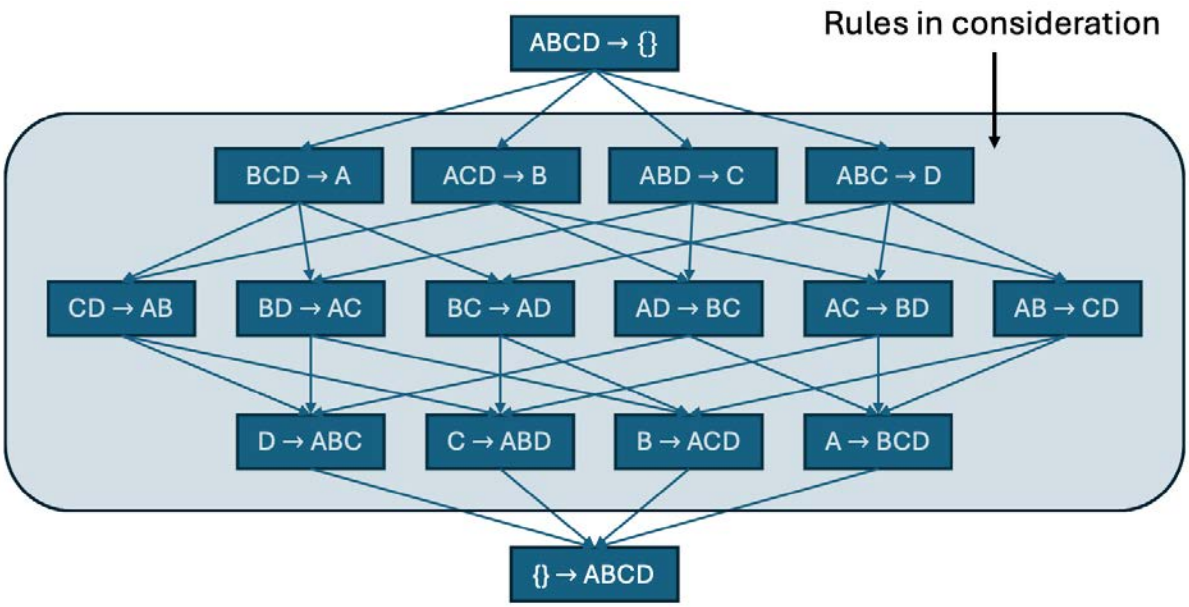

Figure 5: Candidate Rules for a Frequent 4-Itemset.

Once we have the candidate rules, we can determine each rule's Confidence level. The candidate rules whose Confidence level is above the threshold are the association rules of interest.

Now, we need to find a mechanism for reducing the number of candidate rules to evaluate. In general, Confidence does not have the anti-monotone property that Support has. However, Confidence of rules generated from the same itemset has anti-monotone properties. Let us understand through an example.

Suppose we have a frequent itemset {A, B, C, D}, then c({A, B, C} → {D}) ≥ c({A, B} → {C, D}) ≥ c({A} → {B, C, D}).

This is so for the following reasons.

c({A, B, C} → {D}) = $\frac{s(\{A,B,C,D\}}{s(\{A,B,C\}}$

c({A, B} → {C, D}) = $\frac{s(\{A,B,C,D\}}{s(\{A,B\}}$

Now, s({A, B, C}) ≤ s({A, B})

As the numerators are the same, the fraction with the smaller denominator will have the higher value. So, c({A, B, C} → {D}) ≥ c({A, B} → {C, D}).

So, if the confidence of the association rule {B, C, D} → {A} is less than the set threshold, we can discard all the rules below it from the computation. This is shown in the figure below.

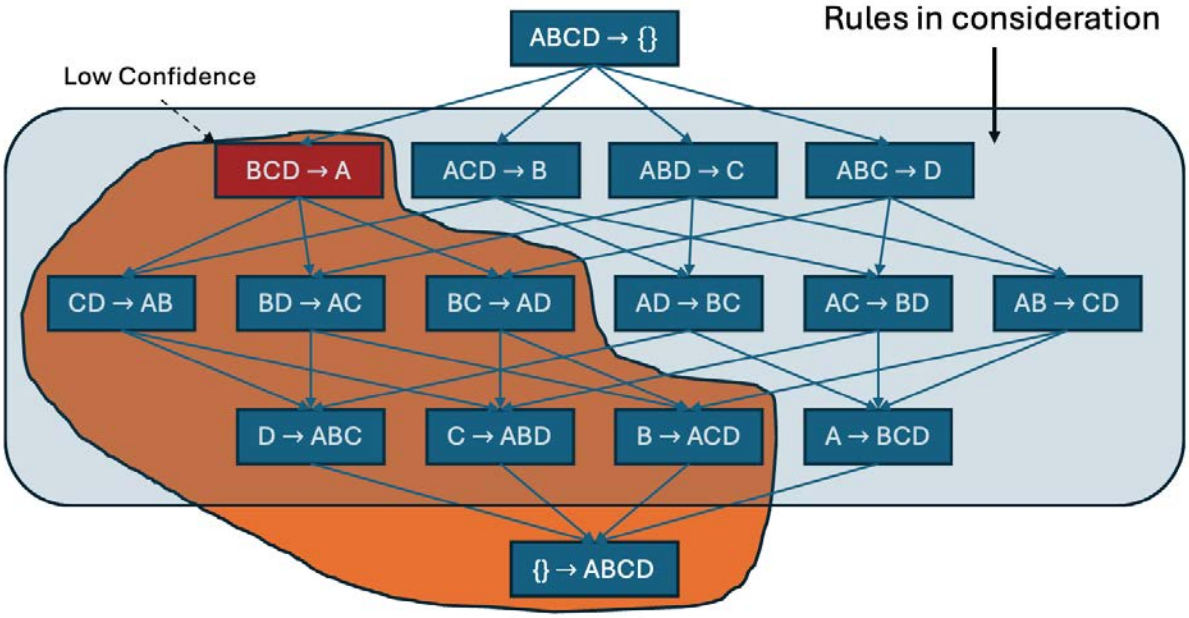

Figure 6: The association rules in the highlighted portion in orange are discarded from further computation.

4.2Pseudo-code for Generating Association Rules

Below is the pseudo-code for implementing the strategy discussed in 4.1.

Input: Frequent itemsets along with their support values and the minimum confidence threshold.
Output: Association rules that satisfy the minimum confidence threshold.

1. **Initialize**
    - Create an empty list called Rules to store valid association rules.
2. **For Each Frequent Itemset in the List of Frequent Itemsets**
    - Generate all non-empty subsets of the frequent itemset.
    - For each subset, do the following:
        - Compute the Consequent by subtracting the subset from the frequent itemset.
        - If the Consequent is not empty, calculate the Confidence using the formula:
            - Confidence equals the support of the frequent itemset divided by the support of the subset.
        - If the Confidence is greater than or equal to the minimum confidence threshold, add the rule (subset implies Consequent) to the Rules list.
3. **Optimize Candidate Rules**
    - Use the anti-monotone property of Confidence to optimize:

- If the Confidence of a rule is below the minimum confidence threshold, discard all rules where the subset is a proper subset of the current subset being evaluated.

4. **Return**
   - Output the list of valid association rules stored in Rules.

4.3Function to Generate Association Rules

Below is the Python implementation for the pseudo-code provided in 4.2.

```
from itertools import combinations

# Function to generate candidate itemsets
def generate_candidates(L, k):
    return set([frozenset(sorted(set(i) | set(j))) for i in L for j in L if len(set(i) |
 set(j)) == k])

# Apriori Algorithm to find frequent itemsets
def apriori(transactions, min_support):
    items = set(item for transaction in transactions for item in transaction)
    itemset = set(frozenset([item]) for item in items)
    L = [set()]
    k = 1
    support_data = {}

    while itemset:
        # Count support for each itemset
        counts = {i: sum([i.issubset(transaction) for transaction in transactions]) for
i in itemset}
        # Retain only itemsets meeting minimum support
        L_k = set(i for i in counts if counts[i] >= min_support)
        L.append(L_k)
        # Store support data
        for i in L_k:
            support_data[i] = counts[i] / len(transactions)
        k += 1
        # Generate candidates for the next iteration
        itemset = generate_candidates(L_k, k)

    return L[1:], support_data

# Function to generate association rules from frequent itemsets
def generate_association_rules(frequent_itemsets, support_data, min_confidence):
    rules = []

    for itemset in frequent_itemsets:
        subsets = [set(subset) for i in range(1, len(itemset)) for subset in combination
s(itemset, i)]
        for subset in subsets:
            consequent = itemset.difference(subset)
            if consequent:
                confidence = support_data[frozenset(itemset)] / support_data[frozenset(s
ubset)]
                if confidence >= min_confidence:
                    rules.append((subset, consequent, confidence))

    return rules

# Dataset from Table 1
transactions = [
```

```
    {'Bread', 'Butter', 'Milk'},
    {'Bread', 'Milk'},
    {'Butter', 'Milk'},
    {'Bread', 'Butter'},
    {'Bread', 'Cheese', 'Butter'},
    {'Milk', 'Cheese'},
    {'Cheese'},
    {'Bread', 'Milk', 'Butter'}
]

# Parameters
min_support = 4  # 50% of 8 transactions
min_confidence = 0.6

# Run Apriori Algorithm
frequent_itemsets, support_data = apriori(transactions, min_support)

# Generate Association Rules
rules = generate_association_rules(
    [itemset for level in frequent_itemsets for itemset in level],
    support_data,
    min_confidence
)

# Output results
print("Frequent Itemsets:")
for k, itemsets in enumerate(frequent_itemsets, 1):
    print(f"{k}-itemsets:")
    for itemset in itemsets:
        print(f"  {set(itemset)}")

print("\nAssociation Rules:")
for antecedent, consequent, confidence in rules:
    print(f"If {set(antecedent)} then {set(consequent)}, Confidence: {confidence:.2f}")

Frequent Itemsets:
1-itemsets:
  {'Milk'}
  {'Bread'}
  {'Butter'}
2-itemsets:
  {'Butter', 'Bread'}

Association Rules:
If {'Butter'} then {'Bread'}, Confidence: 0.80
If {'Bread'} then {'Butter'}, Confidence: 0.80
```

Code 3: AssociationRulesGenerator.ipynb (https://drive.google.com/file/d/1XaOngFV3d0WXIllmLNCFEQYzEKRs0xhf/view?usp=sharing)

4.4Conclusion

In this chapter, we explored the complex journey from frequent itemsets to practical association rules. By using techniques like confidence-based filtering and taking advantage of properties like anti-monotonicity, we simplified the process of discovering rules with significance. The methods mentioned illustrate the effectiveness of association rule mining in revealing hidden relationships between variables in large datasets, leading to applications that improve customer engagement, optimise processes, and predict trends. Now that we know all of these and how they relate to one another, we can finally move on to the next

chapter, where we will apply these concepts and create a recommendation system for a Farmer’s Market, a tangible real-world solution.

Chapter 5: Suggesting Items for Farmer's Market

A Farmer's Market is where farmers come together to sell products directly to consumers, usually off the back of their trucks. Many markets include fish, meat, dairy products, and fresh fruits and vegetables. This concept is gaining popularity as consumers get fresh food products directly from farms at relatively lower prices. At the same time, it helps small farmers earn more. The idea has been practised in many countries, including the USA, Italy, and Thailand, for over 25 years. Though informal Farmer Markets exist in many countries, organised Farmer Markets are restricted to some developed countries.

The concept of Farmer's Markets has recently gained popularity in developing and underdeveloped countries. The most recent example is India, where new Farm Laws were introduced so that farmers could sell their produce directly to consumers. Though the laws were withdrawn after being introduced, the concept is useful because it connects farmers directly with consumers without any middlemen.

With this concept likely to become prevalent in the future, it is essential to establish mechanisms for identifying details, including where Farmer's Markets could be established, what products could be sold in them, what margins the farmers and organisers could make, and how beneficial such Markets could be for consumers.

In this chapter, we will establish a mechanism for selecting products for a new Farmer's Market based on the data collected from established ones.

The data is for the USA. The chapter will demonstrate how to analyse the available data to obtain conclusive answers for the scope. Provided appropriate data is available, the mechanism is expected to be replicated for other countries.

5.1 About the Data

The dataset for this project is obtained from John Snow Labs's US Farmers Market Location dataset. It contains 8,677 rows and 59 columns.

The Python code to load the data is provided below.

```
import pandas as pd
import gdown

# Google Drive file ID
file_id = "1rO1UoiyJi8Jn4Hm1QaVPpuiDTJyvBVae"
url = f"https://drive.google.com/uc?id={file_id}"
```

```
# Download the file using gdown
output_file = "farmers_market_data.csv"
gdown.download(url, output_file, quiet=False)

# Load the dataset into a Pandas DataFrame
df = pd.read_csv(output_file)
```
Code 4: Code to load Farmer Market Dataset.

The list of columns is provided below.

```
'FMID', 'MarketName', 'Website', 'Facebook', 'Twitter', 'Youtube', 'OtherMedia', 'street', 'city',
'County', 'State', 'zip', 'Season1Date', 'Season1Time', 'Season2Date', 'Season2Time', 'Season3Date
', 'Season3Time', 'Season4Date', 'Season4Time', 'x', 'y', 'Location', 'Credit', 'WIC', 'WICcash', '
SFMNP', 'SNAP', 'Organic', 'Bakedgoods', 'Cheese', 'Crafts', 'Flowers', 'Eggs', 'Seafood', 'Herbs',
 'Vegetables', 'Honey', 'Jams', 'Maple', 'Meat', 'Nursery', 'Nuts', 'Plants', 'Poultry', 'Prepared
', 'Soap', 'Trees', 'Wine', 'Coffee', 'Beans', 'Fruits', 'Grains', 'Juices', 'Mushrooms', 'PetFood
', 'Tofu', 'WildHarvested', 'updateTime'
```
Table 6: List of columns in the original dataset

Our objective is to find possible products to sell in the new farmer's markets. So, we eliminate all the data elements that will not lead to this conclusion.

First, we remove the columns containing social media information. In any case, most of the information in these columns is unavailable in the original dataset.

```
'Website', 'Facebook', 'Twitter', 'Youtube','OtherMedia'
```
Table 7: Columns with information regarding social media in the original dataset

The column **updateTime** has no relevance in our analysis, thus removing it.

The original dataset contains columns that indicate the year when the farmer markets were active. This is valuable information. However, since we will not need to provide this level of detail, we will remove these columns.

```
'Season1Date','Season1Time', 'Season2Date', 'Season2Time', 'Season3Date', 'Season3Time', 'Season4Da
te', 'Season4Time'
```
Table 8: Columns with information on when the farmer markets are active during the year in the original dataset.

There is a column named **Location,** which contains information about the whereabouts of the farmer markets. Since we are concerned about the products, we will not use this data. However, in real situations, this data is essential, as the similarity of locations is critical in suggesting products to sell.

The addresses of existing Farmer Markets are essential to finding appropriate locations for new Ones. However, as we are only suggesting products, we will not use the columns containing data regarding Farmer Market addresses.

```
'city', 'County', 'State', 'street', 'zip', 'x', 'y',
```
Table 9: Columns with information on farmer markets addresses in the original dataset.

The columns in the original dataset contain information on the mode of money transactions at the Farmer Markets. As this is not the interest of our study, we will drop these columns.

```
'Credit', 'WIC', 'WICcash', 'SFMNP', 'SNAP'
```
Table 10: List of columns containing information on the mode of money transactions in the original dataset.

So, the final list of columns to use for our analysis is given below.

```
'FMID', 'MarketName', 'Organic', 'Bakedgoods', 'Cheese', 'Crafts', 'Flowers', 'Eggs', 'Seafood', 'H
erbs', 'Vegetables', 'Honey', 'Jams', 'Maple', 'Meat', 'Nursery', 'Nuts', 'Plants', 'Poultry', 'Pre
pared', 'Soap', 'Trees', 'Wine', 'Coffee', 'Beans', 'Fruits', 'Grains', 'Juices', 'Mushrooms', 'Pet
Food', 'Tofu', 'WildHarvested'
```

Table 11: List of columns in the original dataset to be used for the analysis.

5.1.1 Imputing Missing Values

After dropping the columns which are not needed, the list of missing values in the remaining columns as given below.

```
FMID                0
MarketName          0
Organic             0
Bakedgoods       3039
Cheese           3039
Crafts           3039
Flowers          3039
Eggs             3039
Seafood          3039
Herbs            3039
Vegetables       3039
Honey            3039
Jams             3039
Maple            3039
Meat             3039
Nursery          3039
Nuts             3039
Plants           3039
Poultry          3039
Prepared         3039
Soap             3039
Trees            3039
Wine             3039
Coffee           3039
Beans            3039
Fruits           3039
Grains           3039
Juices           3039
Mushrooms        3039
PetFood          3039
Tofu             3039
WildHarvested    3039
```

Table 12: Missing values in the columns to be used for the analysis in the original dataset.

5.1.1.1 Imputing Missing Values for Product Data

The columns containing the data regarding the products sold in the farmer markets have a value of **Y** if the product is sold in a farmer's market and **N** if the product is not sold in the farmer's market. If the value is missing for a particular cell (intersection of product column and farmer’s market row), we will consider the value as **N** (i.e., the product is not sold in that farmer’s market). However, if none of the products have a value of **Y** for a farmer's market, we will not consider that row for our analysis.

The code below imputes the product data.

```
# List of product columns
product_columns = [
    'Organic', 'Bakedgoods', 'Cheese', 'Crafts', 'Flowers', 'Eggs', 'Seafood',
    'Herbs', 'Vegetables', 'Honey', 'Jams', 'Maple', 'Meat', 'Nursery', 'Nuts',
    'Plants', 'Poultry', 'Prepared', 'Soap', 'Trees', 'Wine', 'Coffee', 'Beans',
    'Fruits', 'Grains', 'Juices', 'Mushrooms', 'PetFood', 'Tofu', 'WildHarvested'
]

# Impute missing values in product columns
rows_with_no_products = []

# Iterate through each row in the dataset
for index, row in df.iterrows():
    # Check if all product columns are null or have no 'Y' value
    if row[product_columns].isnull().all() or not (row[product_columns] == 'Y').any():
        rows_with_no_products.append(index)  # Mark rows with no products for removal
    else:
        # Impute missing values in product columns with 'N'
        df.loc[index, product_columns] = row[product_columns].fillna('N')

# Drop rows where no products are sold
df = df.drop(rows_with_no_products).reset_index(drop=True)

# Save the cleaned dataset for further analysis
cleaned_file_path = 'cleaned_farmers_market.csv'
df.to_csv(cleaned_file_path, index=False)
```

Code 5: Code for imputing product data.

After imputing the product data, we have **5,669** rows containing all information about products sold by the existing Farmer Markets.

5.2 Modelling for Creating a Recommendation System

We will create a model to suggest what products to sell in the new Farmer's Market. To indicate the product mix, we will conduct a **Market Basket Analysis**.

First, we create the Frequent Item Set using the **Apriori** algorithm. The minimum support is set to 0.5 so that the model can be built using laptops. At this setting, 319 Frequent Itemsets are extracted. With larger machines, the minimum support can be set to a smaller number to get more frequent itemsets.

```
dfProducts = df.drop(['FMID', 'MarketName'], axis = 1, inplace = False)

# Convert all 'Y' to True and everything else to False in the dataframe
dfProducts = dfProducts.applymap(lambda x: True if x == 'Y' else False)

# Ensure all columns are explicitly boolean type
dfProducts = dfProducts.astype(bool)

!pip install mlxtend

from mlxtend.frequent_patterns import apriori
import pickle

frequentItemsets = apriori(dfProducts, min_support = 0.5, use_colnames = True, low_memor
```

```
y = True)
with open('./FrequentItemsets.obj', 'wb') as fh:
       pickle.dump(frequentItemsets, fh)

frequentItemsets
```

	support	itemsets
0	0.878285	(Bakedgoods)
1	0.502734	(Cheese)
2	0.610513	(Crafts)
3	0.684424	(Flowers)
4	0.744399	(Eggs)
...	...	...
314	0.508203	(Jams, Vegetables, Honey, Soap, Herbs)
315	0.519316	(Jams, Vegetables, Honey, Herbs, Fruits)
316	0.520903	(Jams, Bakedgoods, Eggs, Vegetables, Honey, He...
317	0.509084	(Jams, Bakedgoods, Eggs, Vegetables, Honey, Fr...
318	0.503263	(Jams, Bakedgoods, Vegetables, Honey, Herbs, F...

319 rows × 2 columns

Code 6: Code to find the Frequent Itemsets.

Once the Frequent Item Set is created, the Association Rules are formed using the **Confidence** criterion. The minimum threshold was set to 0.5 so that the model could be built using laptops. The model returned 3,302 Association Rules.

```
from mlxtend.frequent_patterns import association_rules
import pickle

# Ensure 'frequentItemsets' DataFrame has the required 'support' and 'itemsets' columns
frequentItemsets['itemsets'] = frequentItemsets['itemsets'].apply(lambda x: frozenset
(x))  # Convert 'itemsets' to frozenset

# Generate association rules using confidence as the metric
rules = association_rules(frequentItemsets, metric='confidence', min_threshold=0.5, num_
itemsets=dfProducts.shape[0])

# Save the rules to a file
with open('./Rules.obj', 'wb') as fh:
    pickle.dump(rules, fh)

# Sort rules by confidence in descending order and display
sorted_rules = rules.sort_values(by='confidence', ascending=False)
print(sorted_rules)

                            antecedents                        consequents  \
2127  (Flowers, Herbs, Jams, Bakedgoods)                       (Vegetables)
895         (Flowers, Herbs, Bakedgoods)                       (Vegetables)
1623              (Eggs, Herbs, Flowers)                       (Vegetables)
1679              (Flowers, Herbs, Jams)                       (Vegetables)
2908       (Flowers, Honey, Jams, Herbs)                       (Vegetables)
...                                  ...                                ...
2423                        (Vegetables)     (Eggs, Honey, Soap, Bakedgoods)
2602                        (Vegetables)  (Plants, Honey, Herbs, Bakedgoods)
1756                        (Vegetables)                 (Eggs, Herbs, Soap)
577                         (Vegetables)                   (Eggs, Prepared)
1532                        (Vegetables)        (Plants, Fruits, Bakedgoods)
```

```
      antecedent support  consequent support   support  confidence      lift  \
2127            0.541365            0.963838  0.533251    0.985011  1.021967
895             0.585465            0.963838  0.576469    0.984634  1.021576
1623            0.517199            0.963838  0.509084    0.984311  1.021241
1679            0.555830            0.963838  0.547010    0.984132  1.021055
2908            0.516493            0.963838  0.508203    0.983948  1.020864
...                  ...                 ...       ...         ...       ...
2423            0.963838            0.512260  0.502381    0.521230  1.017511
2602            0.963838            0.511201  0.502381    0.521230  1.019618
1756            0.963838            0.510496  0.502029    0.520864  1.020310
577             0.963838            0.511378  0.501499    0.520315  1.017477
1532            0.963838            0.509790  0.500794    0.519583  1.019209

      representativity  leverage  conviction  zhangs_metric   jaccard  \
2127               1.0  0.011462    2.412606       0.046868  0.548639
895                1.0  0.012175    2.353339       0.050949  0.592566
1623               1.0  0.010588    2.304908       0.043080  0.523775
1679               1.0  0.011280    2.278903       0.046425  0.562387
2908               1.0  0.010387    2.252790       0.042270  0.522773
...                ...       ...         ...            ...       ...
2423               1.0  0.008646    1.018736       0.475911  0.515942
2602               1.0  0.009666    1.020947       0.532064  0.516503
1756               1.0  0.009993    1.021639       0.550465  0.516328
577                1.0  0.008614    1.018631       0.474989  0.515036
1532               1.0  0.009439    1.020384       0.521192  0.514778

      certainty  kulczynski
2127   0.585510    0.769135
895    0.575072    0.791365
1623   0.566143    0.756248
1679   0.561192    0.775832
2908   0.556106    0.755609
...         ...         ...
2423   0.018391    0.750973
2602   0.020517    0.751988
1756   0.021181    0.752139
577    0.018290    0.750499
1532   0.019976    0.750968

[3302 rows x 14 columns]
```

Code 7: Code to generate the Association Rules.

For example, using the Association Rules, let us find the recommended products to include in the new Farmer's Market if "Baked Goods" is considered one of the products to include.

```
import matplotlib.pyplot as plt

pProductList = ['Bakedgoods']
inputProductList = frozenset(pProductList)
rules[rules["antecedents"] == inputProductList].groupby(['antecedents', 'consequents'])
[['confidence']].max().sort_values(ascending=False,
                 by='confidence').head(10).plot(kind='bar').invert_xaxis()
plt.title('Top Products that are likely to be sold with ' + str(pProductList))
plt.show()
```

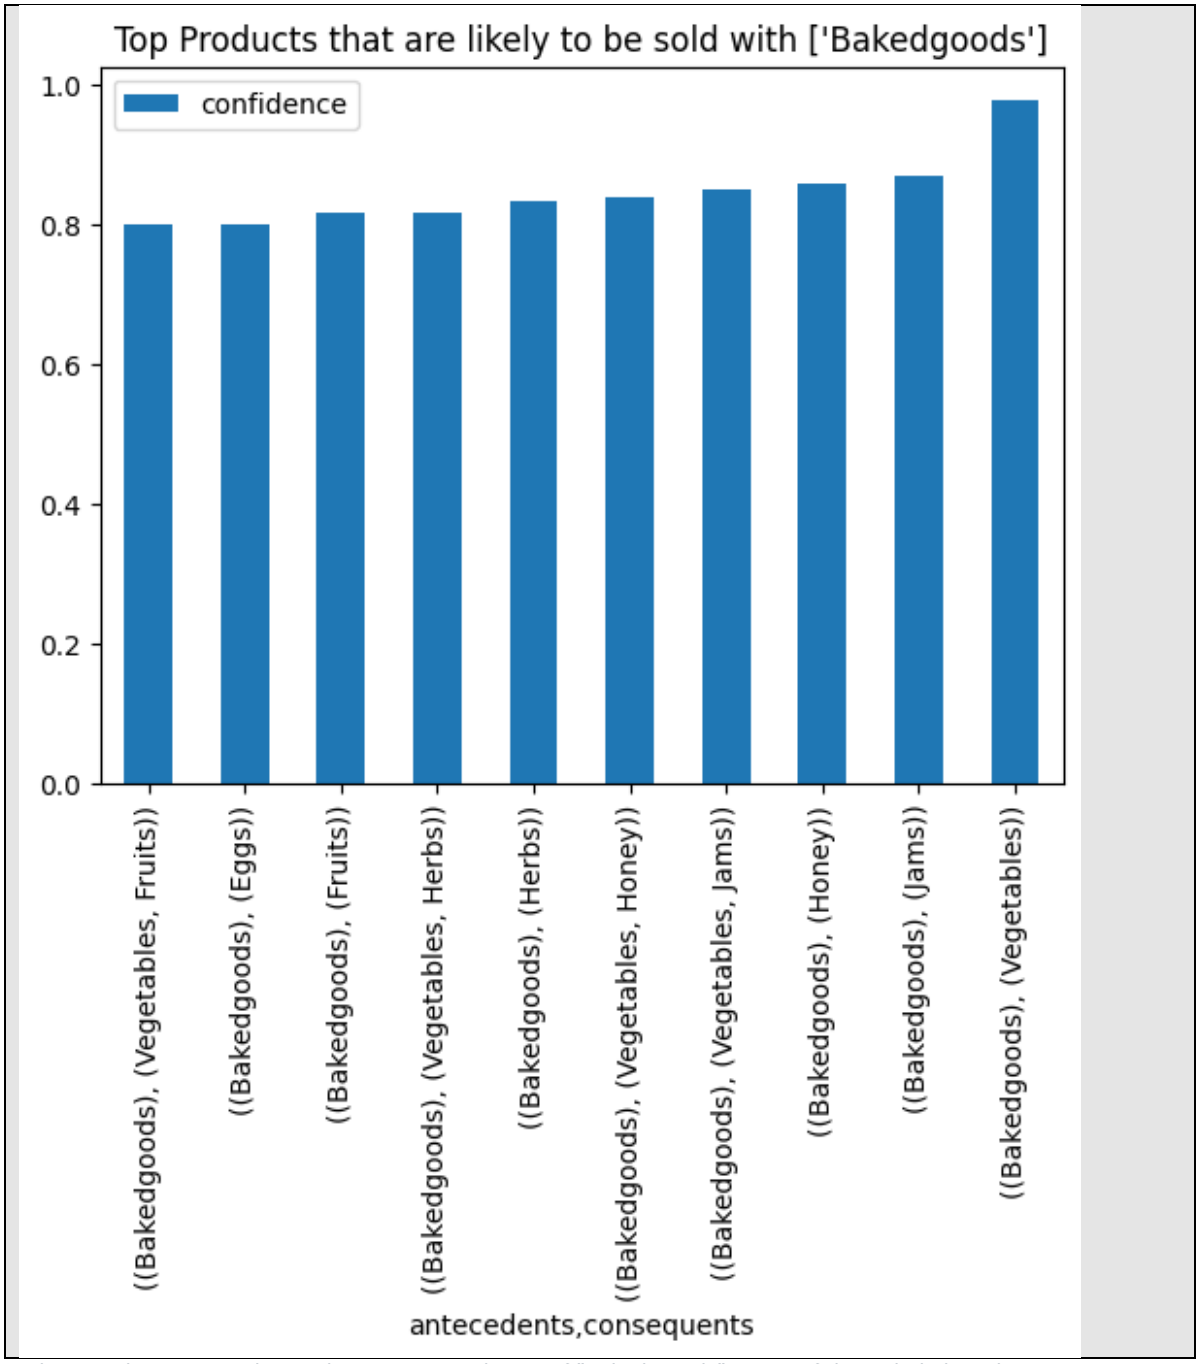

Code 8: Code to create the product recommendations if "Baked Goods" is one of the included products.

Now, we visualise the top 100 products in terms of Confidence.

```
!pip install adjustText

import networkx as nx
import matplotlib.pyplot as plt
from adjustText import adjust_text

# Apply a modern and elegant plot style
plt.style.use('Solarize_Light2')

# Clean and preprocess node labels to remove 'frozenset'
```

```python
rules['antecedents'] = rules['antecedents'].apply(lambda x: ', '.join(list(x)))
rules['consequents'] = rules['consequents'].apply(lambda x: ', '.join(list(x)))

# Create the graph from the top 100 association rules sorted by confidence
edges = nx.from_pandas_edgelist(
    rules.sort_values(ascending=False, by='confidence').head(100),
    source='antecedents',
    target='consequents',
    edge_attr=None
)

# Set up the figure
plt.figure(figsize=(40, 30))
plt.suptitle(
    'Top 100 Products in terms of Confidence',
    fontsize=50,
    fontweight='bold'
)

# Use a visually appealing layout
pos = nx.spring_layout(edges, seed=42)

# Draw the nodes with a pleasant color and transparency
nx.draw_networkx_nodes(
    edges, pos,
    node_size=2000,
    alpha=0.8,
    node_color='orange'
)

# Draw the edges with subtle visibility
nx.draw_networkx_edges(
    edges, pos,
    width=3,
    alpha=0.3,
    edge_color='teal'
)

# Add text labels with dynamic adjustment for clarity
texts = []
for node, (x, y) in pos.items():
    texts.append(plt.text(x, y, node, fontsize=18, fontweight='bold', color='darkblue'))

# Dynamically adjust text positions to prevent overlap
adjust_text(texts, arrowprops=dict(arrowstyle='-', color='gray', alpha=0.5))

# Remove gridlines and axes
plt.grid(False)
plt.axis('off')

# Save and display the plot
plt.tight_layout()
plt.savefig('./AssociationRulesConfidence.png', dpi=300)
plt.show()
```

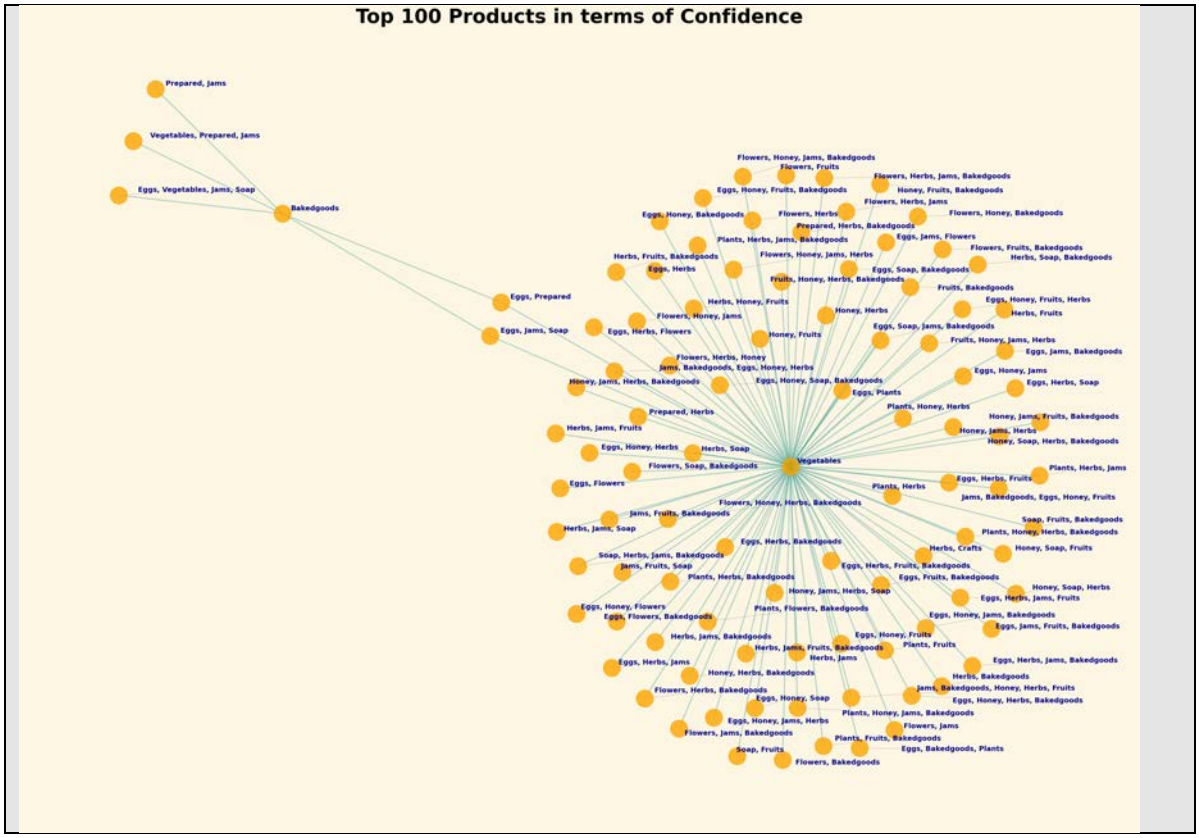

Code 9: Code to visualise the top 100 products by Confidence.

We also visualise the top 100 products by Conviction.

```
import networkx as nx
import matplotlib.pyplot as plt
from adjustText import adjust_text

# Apply a modern and elegant plot style
plt.style.use('Solarize_Light2')

# Create the graph from the top 100 association rules sorted by conviction
edges = nx.from_pandas_edgelist(
    rules.sort_values(ascending=False, by='conviction').head(100),
    source='antecedents',
    target='consequents',
    edge_attr=None
)

# Set up the figure
plt.figure(figsize=(40, 30))
plt.suptitle(
    'Top 100 Products in terms of Conviction',
    fontsize=50,
    fontweight='bold'
)

# Use a visually appealing layout
pos = nx.spring_layout(edges, seed=42)

# Draw the nodes with a pleasant color and transparency
nx.draw_networkx_nodes(
    edges, pos,
```

```
    node_size=2500,
    alpha=0.8,
    node_color='lightcoral'
)

# Draw the edges with subtle visibility
nx.draw_networkx_edges(
    edges, pos,
    width=4,
    alpha=0.3,
    edge_color='mediumslateblue'
)

# Add labels without "frozenset" and adjust them
texts = []
for node, (x, y) in pos.items():
    label = ', '.join(node) if isinstance(node, frozenset) else str(node)  # Convert fro
zenset to a readable string
    texts.append(plt.text(x, y, s=label, fontsize=22, fontweight='bold', color='darkgree
n'))

# Dynamically adjust label positions to avoid overlaps
adjust_text(texts, arrowprops=dict(arrowstyle="-", color='gray', alpha=0.5))

# Remove gridlines and axes
plt.grid(False)
plt.axis('off')

# Save and display the plot
plt.tight_layout()
plt.savefig('./AssociationRulesConviction.png', dpi=300)
plt.show()
```

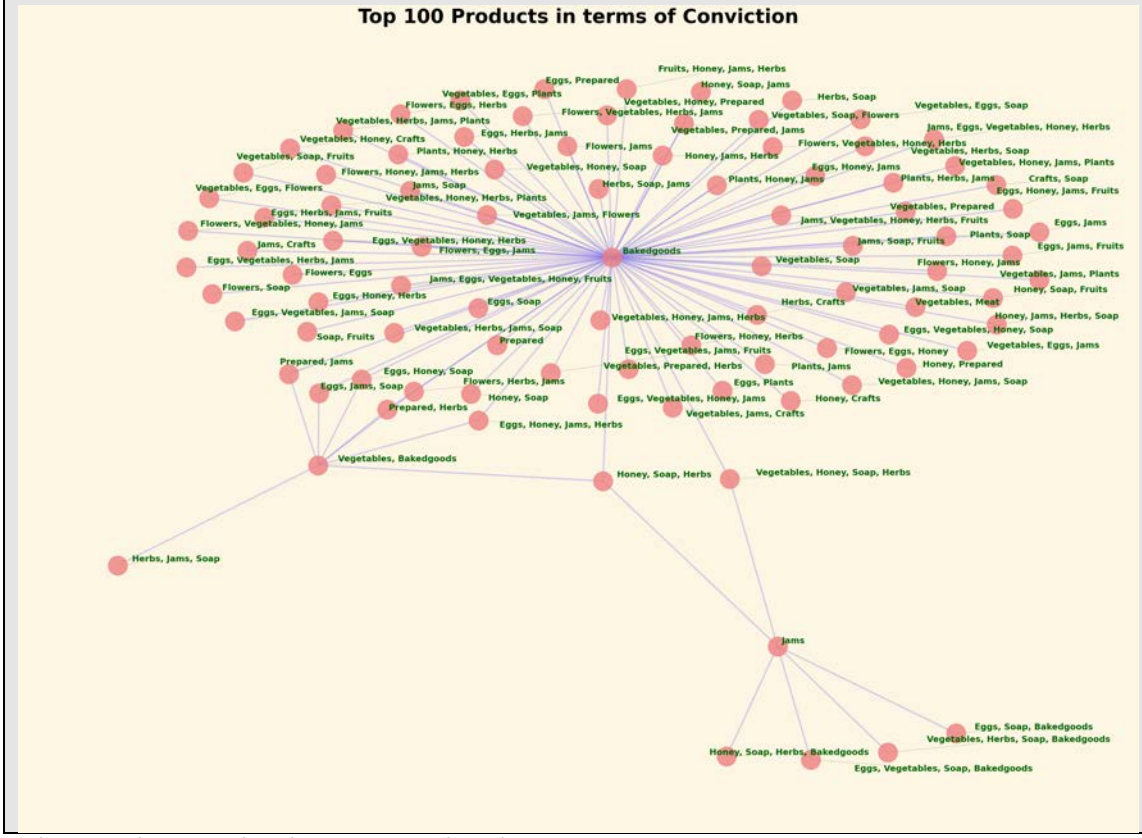

Code 10: Code to visualise the top 100 products by Conviction.

Lastly, we visualise the top 100 products by Certainty.

```
import networkx as nx
import matplotlib.pyplot as plt
from adjustText import adjust_text

# Apply a modern and elegant plot style
plt.style.use('Solarize_Light2')

# Ensure 'certainty' column exists in your rules DataFrame
if 'certainty' not in rules.columns:
    raise ValueError("The 'rules' DataFrame must contain a 'certainty' column.")

# Create the graph from the top 100 association rules sorted by certainty
edges = nx.from_pandas_edgelist(
    rules.sort_values(ascending=False, by='certainty').head(100),
    source='antecedents',
    target='consequents',
    edge_attr=None
)

# Set up the figure
plt.figure(figsize=(40, 30))
plt.suptitle(
    'Top 100 Products in terms of Certainty',
    fontsize=50,
    fontweight='bold'
)

# Use a visually appealing layout
pos = nx.spring_layout(edges, seed=42)

# Draw the nodes with a pleasant color and transparency
nx.draw_networkx_nodes(
    edges, pos,
    node_size=2500,
    alpha=0.8,
    node_color='lightskyblue'
)

# Draw the edges with subtle visibility
nx.draw_networkx_edges(
    edges, pos,
    width=4,
    alpha=0.3,
    edge_color='mediumslateblue'
)

# Add labels without "frozenset" and adjust them
texts = []
for node, (x, y) in pos.items():
    label = ', '.join(node) if isinstance(node, frozenset) else str(node)  # Convert fro
zenset to a readable string
    texts.append(plt.text(x, y, s=label, fontsize=22, fontweight='bold', color='darkgree
n'))

# Dynamically adjust label positions to avoid overlaps
adjust_text(texts, arrowprops=dict(arrowstyle="-", color='gray', alpha=0.5))

# Remove gridlines and axes
plt.grid(False)
plt.axis('off')

# Save and display the plot
```

```
plt.tight_layout()
plt.savefig('./AssociationRulesCertainty.png', dpi=300)
plt.show()
```

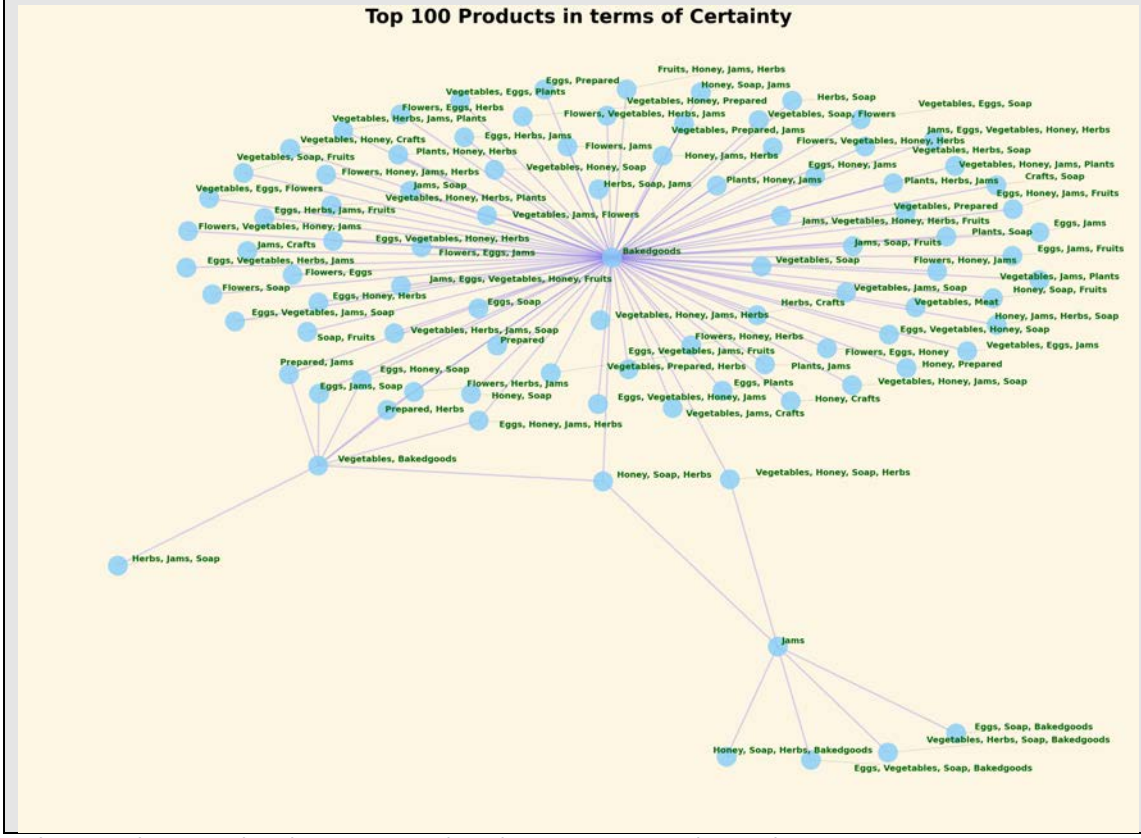

Code 11: Code to visualise the top 100 products by Certainty. Complete Code - https://drive.google.com/file/d/1M9a8Hr1ZdJrpqs1iiCJERc82hKhVtPnH/view?usp=sharing

5.3 Building a Client Application for Recommending Products for Farmer Markets

Now that the model is ready, we built a simple client application using Gradio to illustrate how decision-makers can use such models in real life. The code is provided below.

```
import warnings
warnings.filterwarnings('ignore')

import pandas as pd
import gdown

# Google Drive file ID
file_id = "1rO1UoiyJi8Jn4Hm1QaVPpuiDTJyvBVae"
url = f"https://drive.google.com/uc?id={file_id}"

# Download the file using gdown
output_file = "farmers_market_data.csv"
gdown.download(url, output_file, quiet=False)

# Load the dataset into a Pandas DataFrame
dfOriginal = pd.read_csv(output_file)
```

```
dfOriginal

Downloading...
From: https://drive.google.com/uc?id=1rO1UoiyJi8Jn4Hm1QaVPpuiDTJyvBVae
To: /content/farmers_market_data.csv
100%|██████████| 2.59M/2.59M [00:00<00:00, 110MB/s]

{"variable_name":"dfOriginal","type":"dataframe"}

# Load the Objects
import pickle

# Google Drive file ID
file_id = "10n1i3yi1gg0vBYwHB4MwDILx938wvupi"
url = f"https://drive.google.com/uc?id={file_id}"

# Download the file using gdown
output_file = "FrequentItemsets.obj"
gdown.download(url, output_file, quiet=False)

pickleFrequentItemsets = open ("FrequentItemsets.obj", "rb")
frequentItemsets = pickle.load(pickleFrequentItemsets)

# Google Drive file ID
file_id = "1G1nSSTshotR95V1bYf36cIVv5E4obCb6"
url = f"https://drive.google.com/uc?id={file_id}"

# Download the file using gdown
output_file = "Rules.obj"
gdown.download(url, output_file, quiet=False)

pickleRules = open ("Rules.obj", "rb")
rules = pickle.load(pickleRules)

Downloading...
From: https://drive.google.com/uc?id=10n1i3yi1gg0vBYwHB4MwDILx938wvupi
To: /content/FrequentItemsets.obj
100%|██████████| 6.42k/6.42k [00:00<00:00, 3.87MB/s]
Downloading...
From: https://drive.google.com/uc?id=1G1nSSTshotR95V1bYf36cIVv5E4obCb6
To: /content/Rules.obj
100%|██████████| 366k/366k [00:00<00:00, 83.5MB/s]

lProducts = ['Organic', 'Bakedgoods', 'Cheese', 'Crafts', 'Flowers', \
             'Eggs', 'Seafood', 'Herbs', 'Vegetables', 'Honey', 'Jams', \
             'Maple', 'Meat', 'Nursery', 'Nuts', 'Plants', 'Poultry', 'Prepared', \
             'Soap', 'Trees', 'Wine', 'Coffee', 'Beans', 'Fruits', 'Grains', 'Juices', \
             'Mushrooms', 'PetFood', 'Tofu', 'WildHarvested']
lProducts.sort()

def find_related_products(pProductList):
    if not pProductList:
        return pd.DataFrame(columns=["Product Combinations", "Conviction"]), "No product
s selected."

    # Convert the input list to a frozenset for querying
    inputProductList = frozenset(pProductList)

    # Filter rules for the selected products
    filtered_rules = rules[rules["antecedents"] == inputProductList].sort_values(ascendi
ng=False, by="conviction")

    if filtered_rules.empty:
        return pd.DataFrame(columns=["Product Combinations", "Conviction"]), "No related
```

```
 products found."

    # Format the consequents column by converting frozensets to strings
    filtered_rules["consequents"] = filtered_rules["consequents"].apply(lambda x: ", ".j
oin(x))

    # Prepare the DataFrame with related products and conviction scores
    related_products = filtered_rules[["consequents", "conviction"]].rename(
        columns={"consequents": "Product Combinations", "conviction": "Conviction"}
    )

    # Limit the number of results to the top 10
    toShowNumber = min(10, len(related_products))
    message = f"Showing {toShowNumber} out of {len(related_products)} suggestions."

    return related_products.head(toShowNumber), message

!pip install gradio

import gradio as gr
import pandas as pd

# Ensure the list of products `lProducts` and the function `find_related_products` are d
efined elsewhere in your code.

# Define the dropdown for selecting products
product_selector = gr.Dropdown(
    choices=lProducts,  # List of products available for selection
    label="Select Products",  # Label for the dropdown
    multiselect=True,  # Enable multi-select for selecting multiple products
    info="Select one or more products to find related recommendations."  # Add descripti
ve info
)

# Define the dataframe to display related products and convictions
related_products_display = gr.DataFrame(
    headers=["Product Combinations", "Conviction"],  # Table headers
    label="Related Products",  # Label for the DataFrame
    interactive=False # ,  # Make the DataFrame read-only
#    info="Displays related product combinations and their conviction scores."
)

# Define a label for messages or feedback
message_label = gr.Label(
    label="Message" # ,
#     info="Displays the number of recommendations found or other messages."
)

# Define the inputs and outputs for the Gradio interface
inputs = [product_selector]
outputs = [related_products_display, message_label]

# Create the Gradio interface
app = gr.Interface(
    fn=find_related_products,  # Function to call when inputs change
    inputs=inputs,  # Input components
    outputs=outputs,  # Output components
    title="Related Products Finder",  # Title of the app
    description=(
        "This tool suggests related products based on the selected products. "
        "The suggestions are generated using association rules derived from "
        "frequent itemsets, providing insights into potential product combinations."
    ),  # Description of the app
    theme="compact"  # Use a professional, compact theme for better layout
```

```
)

# Launch the app
if __name__ == "__main__":
    app.launch(share=True)  # Share set to False for local usage
```

Code 12: FarmerMarket-MBA-GUI.ipynb (https://drive.google.com/file/d/168pg1NLUkLdxogBOL0p7Uzb3_U3Ka_ii/view?usp=sharing)

A few screenshots of running the above program are shown below.

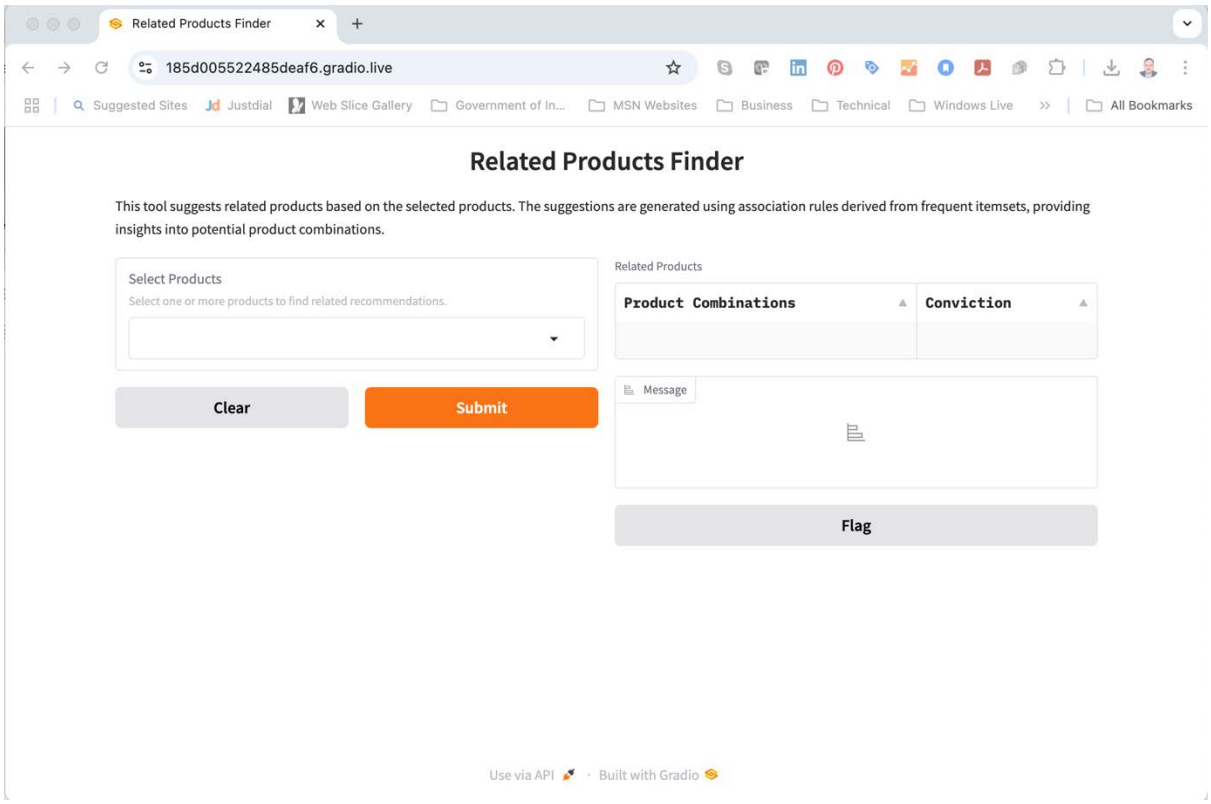

Figure 7: This is the opening screen of the Client GUI.

On the opening screen, one can use the left panel and select products from a drop-down list. These are the products already considered for including the new Farmer Market. One can choose one or more number of products. After selecting the products, the "Submit" button must be pressed to obtain a recommendation on the right-hand panel. This is shown in the figure below.

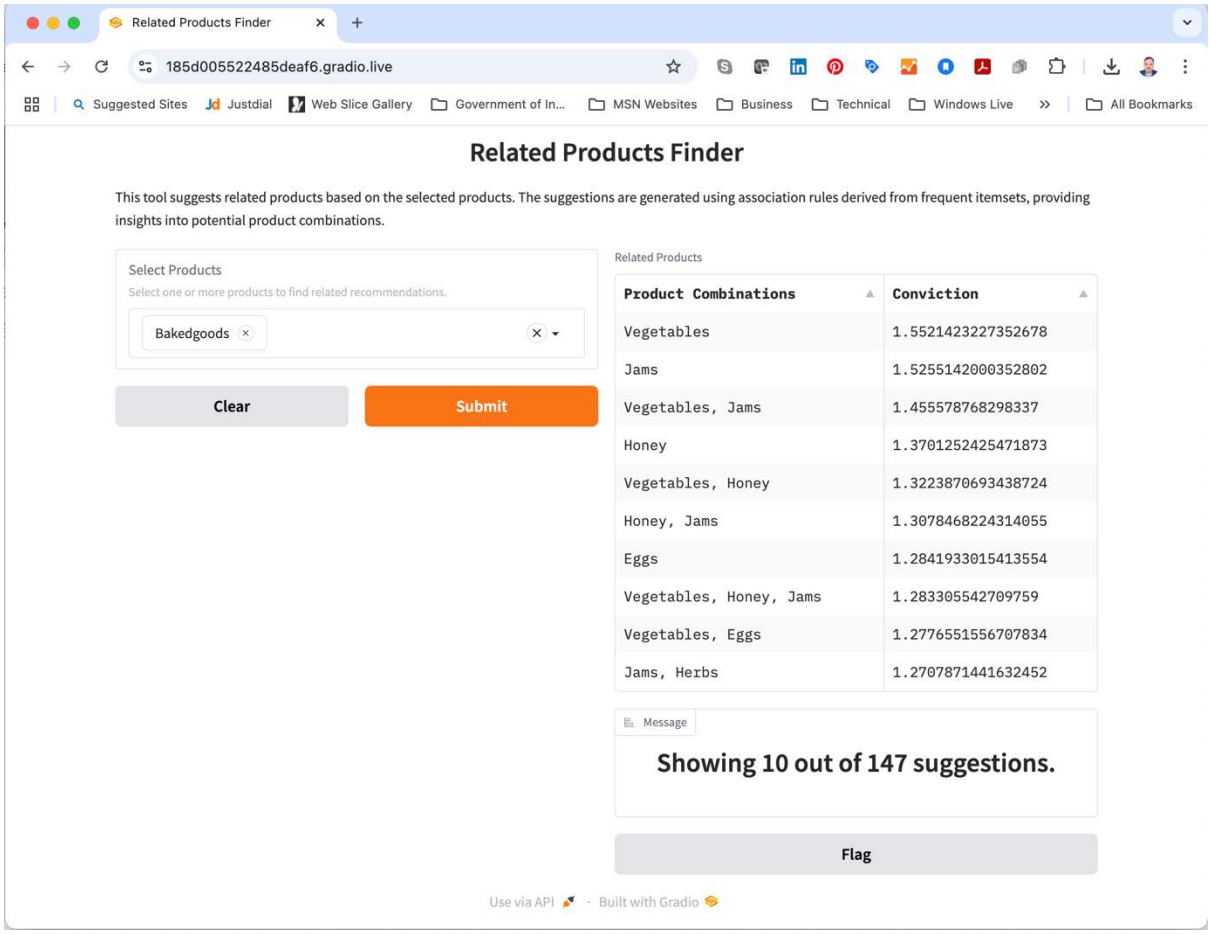

Figure 8: Recommendations are shown on the right-hand panel on selecting products on the left-hand panel and pressing "Submit.

The figure below shows the recommendations when two products are selected on the left-hand panel.

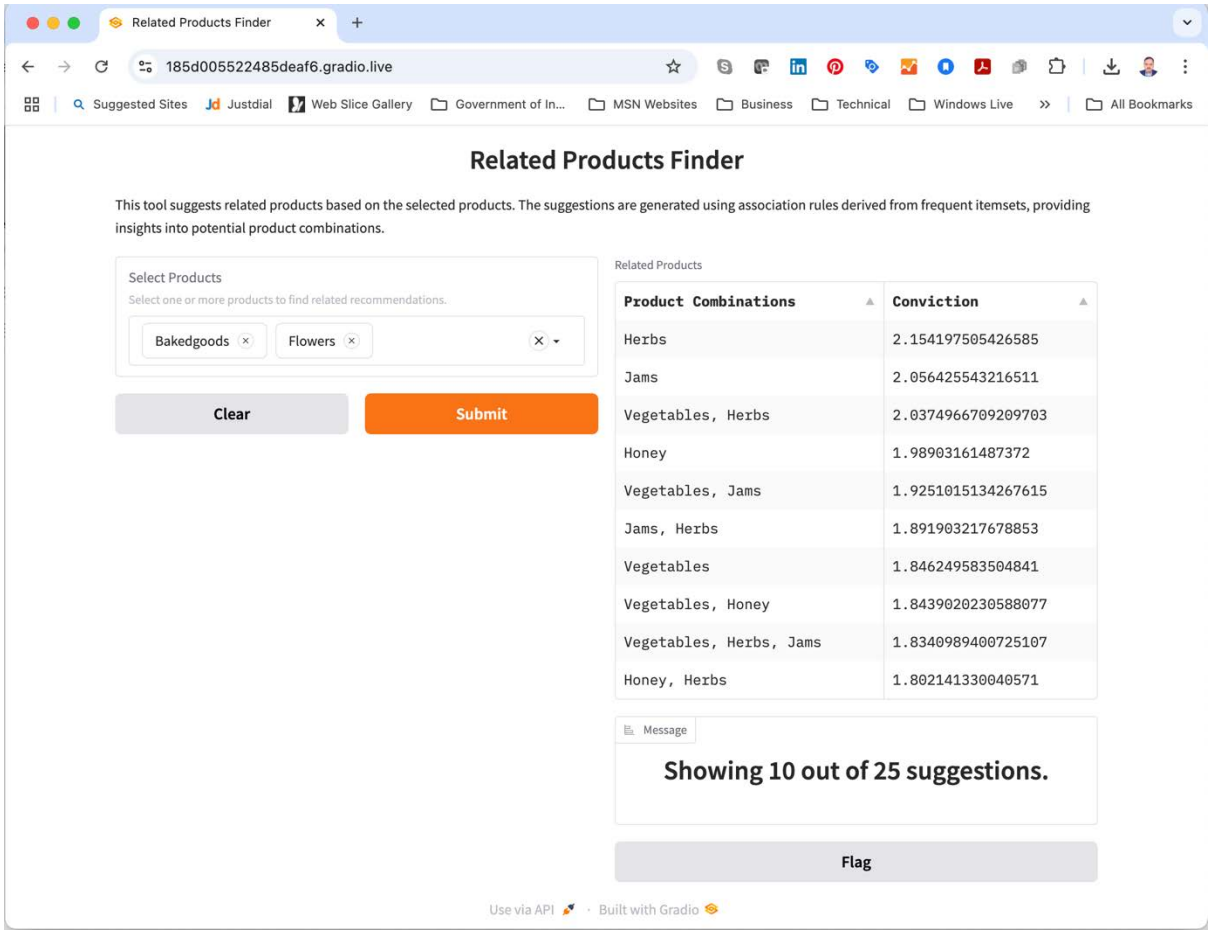

Figure 9: Recommendations generated for the itemset {Bakedgoods, Flowers}.

5.4Conclusion

This book has been an exploration of how data-driven approaches and computational methods can create recommendation systems for diverse domains. By delving into Market Basket Analysis, the Apriori algorithm, association rule generation, and the practical implementation of these concepts, we have laid a comprehensive foundation for developing systems that transform data into actionable insights. While the Farmer's Market served as an illustrative example, the methodologies and principles discussed are universally applicable across industries, making this book a guide for building recommendation systems in various contexts.

The Farmer's Market example demonstrated how data could guide decisions about product combinations, market structures, and consumer preferences. However, the real strength of these techniques lies in their adaptability. From retail and healthcare to education and finance, recommendation systems have become indispensable tools for personalising experiences, optimising operations, and driving growth. This book equips readers with the knowledge to build such systems,

providing the technical steps and the underlying rationale and decision-making frameworks.

The journey through this book emphasised the importance of data preprocessing, algorithmic efficiency, and interpretability in creating reliable and user-centric recommendation systems. By applying these principles, organisations and individuals can unlock the potential of their data, uncovering patterns and relationships that inform strategic decisions. Visualisations, such as those presented in the Farmer's Market example, further enhance the interpretability and accessibility of these insights for stakeholders.

As this book concludes, it invites you to envision recommendation systems beyond what we have explored. Whether it's improving patient care through healthcare recommendations, tailoring educational resources to individual learners, or optimising supply chains in manufacturing, the tools and techniques discussed here are only the beginning. By bridging data science with real-world applications, we aim to inspire innovation and problem-solving across domains.

Thank you for embarking on this journey. May the insights you gain empower you to create intelligent systems that positively impact the industries and communities you serve.

Annexures

List of Code Snippets

Code 1: NumberOfAssociationRules.ipynb
(https://drive.google.com/file/d/1VJyr60gxD7lOe48edePkAK2BRGIO3gTE/view?usp=sharing) 29
Code 2: AprioriAlgorithm.ipynb
(https://drive.google.com/file/d/185F4tXmz6Ao7gu7LQAPFQzkeUNGnYRe1/view?usp=sharing) ..34
Code 3: AssociationRulesGenerator.ipynb
(https://drive.google.com/file/d/1XaOnqFV3d0WXIllmLNCFEQYzEKRs0xhf/view?usp=sharing) ...40
Code 4: Code to load Farmer Market Dataset. 43
Code 5: Code for imputing product data. 45
Code 6: Code to find the Frequent Itemsets. 46
Code 7: Code to generate the Association Rules. 47
Code 8: Code to create the product recommendations if "Baked Goods" is one of the included products. .48
Code 9: Code to visualise the top 100 products by Confidence. 50
Code 10: Code to visualise the top 100 products by Conviction. 51
Code 11: Code to visualise the top 100 products by Certainty. Complete Code -
https://drive.google.com/file/d/1M9a8Hr1ZdJrpqs1iiCJERc82hKhVtPnH/view?usp=sharing 53
Code 12: FarmerMarket-MBA-GUI.ipynb
(https://drive.google.com/file/d/168pg1NLUkLdxogBOL0p7Uzb3_U3Ka_ii/view?usp=sharing) 56

About the Author

Partha Majumdar's leadership in the dynamic realm of software solutions is about technical prowess, strategic insight, and a unique style that blends these qualities with a personal touch. This makes his approach to innovation, efficiency, and business success unique.

His educational journey, which includes a Global Doctorate in Business Administration and specialisations in Computational Data Sciences and Cybersecurity, reflects a commitment to continuous learning.

Majumdar's professional journey has been nothing short of extraordinary. As the Vice President of Software Engineering at J.P. Morgan Chase and Co. in his last role, he has spearheaded impactful initiatives, contributing to the evolution of software development paradigms. He is in the inception phase of starting his firm in the UAE and pursuing his PhD in Computer Science from Kalinga University. His earlier role as the Managing Director of Majumdar Consultancy Private Limited showcased his entrepreneurial spirit, where he successfully nurtured a fledgling business into a success with a global footprint.

Partha Majumdar's diverse talents and expertise are not confined to a single domain. His proficiency in software development, predictive modelling, descriptive data analysis, and Agile Project Management is a testament to his versatility, adaptability, and ability to deliver innovative solutions across various contexts.

Beyond his roles in corporate leadership, Majumdar has been recognised with numerous awards, including the Excellence Award, Gratitude Award, Merit Award, and Best IT Manager, which underline his impact and leadership in the field. His commitment to excellence is further demonstrated through a comprehensive list of professional upskilling, covering project management, IT service management, and specialised areas like data science and cloud computing.

Majumdar's publications and patent attempts showcase a commitment to advancing the field. He has published thirty-seven books on academia and knowledge dissemination, and his upcoming books and publications illustrate his dedication to sharing knowledge and insights.

In conclusion, Partha Majumdar's career is a testament to his multifaceted expertise, from spearheading successful ventures to influencing software development paradigms at industry giants. His unwavering commitment to innovation, coupled with a rich educational background and many certificates, positions him as a distinguished leader poised to continue making impactful contributions in the ever-evolving landscape of software solutions.

Books by the Author

Mastering Classification Algorithms for Machine Learning

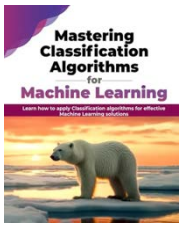

This book explores the fundamentals of machine learning through classification algorithms, which are crucial for categorising inputs in applications like spam detection and fraud prevention. It covers problem-solving in machine learning, delving into Naïve Bayes, Logistic Regression (including the sigmoid function), Decision Trees (focusing on Gini Factor and Entropy), Random Forest, and Boosting techniques, providing a comprehensive guide to mastering classification challenges.

Link in Amazon Store: https://www.amazon.in/dp/935551851X

Neural Networks for Engineers

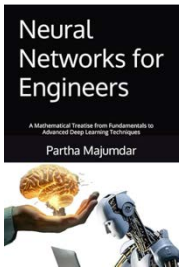

This book offers a comprehensive guide to neural networks, from foundational concepts to advanced deep learning techniques. It bridges the gap between theory and application, covering Perceptrons, CNNs, RNNs, and more, with hands-on examples using Scikit-Learn, Keras, and PyTorch. Designed for engineers and practitioners, it demystifies neural network computations. It provides practical insights into deploying sophisticated models in real-world scenarios, making it an essential resource for mastering neural networks.

Link in Amazon Store: https://www.amazon.in/dp/B0DFVPV8YH

Time Series Analysis for Engineers

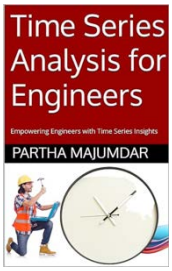

This book is a comprehensive resource for professionals seeking to understand and analyse temporal data. This book bridges theory and practice, covering foundational concepts and statistical methods. With practical Python and R examples, it equips readers to tackle real-world challenges in stock markets, weather patterns, and industrial processes. Whether for forecasting, decision-making, or innovation, this book is an essential guide for engineers and data practitioners.

Link in Amazon Store: https://www.amazon.com/dp/B0DQJYPCRS

Machine Learning for Managers

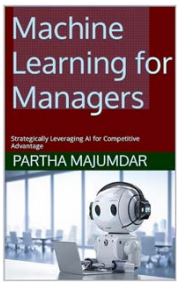

This book is a comprehensive guide tailored for leaders aiming to leverage machine learning (ML) within their organisations. It simplifies ML concepts, emphasising strategic applications over technical complexity. The book covers integrating ML into business practices, ethical data use, and real-world industry applications, showcasing ML's role in enhancing operations and innovation. It also provides insights on team building in the ML era, promoting cross-disciplinary collaboration for effective ML adoption. This book is a strategic roadmap for managers to harness ML, driving informed decision-making and positioning their organisations for future success in an AI-driven landscape.

Link in Amazon Store: https://www.amazon.in/dp/B0CZ5XTQ1L

Deep Learning for Managers

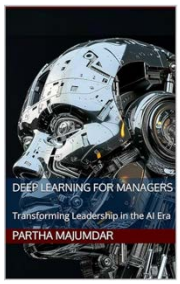

This book is a pivotal guide for modern leaders navigating the AI revolution. It demystifies deep learning, making it accessible to managers without requiring deep technical knowledge. This book equips leaders with the insights to harness AI effectively, covering everything from the basics of artificial neural networks to the ethical considerations of AI deployment. It's an indispensable resource for any leader aiming to leverage deep learning as a strategic asset in today's rapidly evolving business landscape.

Link in Amazon Store: https://www.amazon.in/dp/B0CWDPWSN8

Generative AI for Managers

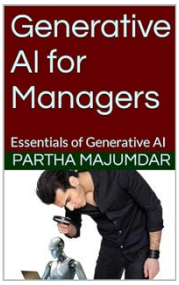

This book is a cutting-edge guide that demystifies Generative AI for business leaders eager to harness this technology for growth and innovation. It delves into how Generative AI can revolutionise aspects of business, from enhancing customer experiences to optimising operations and driving strategic decision-making. The book provides a wealth of practical applications, showcases how mundane tasks can be automated for efficiency, and presents strategies for fostering a culture of innovation through AI. Additionally, it offers guidance on the ethical implementation of AI technologies, ensuring they complement and augment human capabilities within the organisational framework, thereby paving the way for a future rich in opportunities and advancements.

Link in Amazon Store: https://www.amazon.in/dp/B0CXYBFJHD

ChatGPT AI for Managers

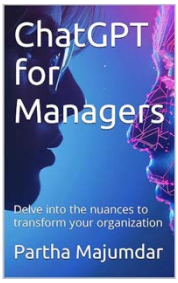

This book is a vital resource for leaders navigating the AI revolution, focusing on integrating Generative AI, like ChatGPT, in enhancing managerial functions and team dynamics. It provides practical insights into leveraging ChatGPT to streamline tasks, bolster decision-making, and encourage innovative thinking within teams. This guide transcends theoretical knowledge, offering actionable strategies for managers to complement their skills with AI, thereby elevating their leadership effectiveness. Through real-world applications and expert advice, readers will learn to harmonise traditional management with AI advancements, ensuring they remain at the forefront of the evolving business environment.

Link in Amazon Store: https://www.amazon.in/dp/B0CY8L4CQ9

Data Lakes for Managers

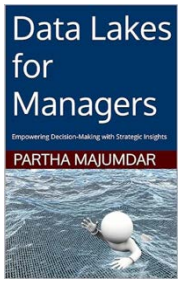

This book is a guide tailored for managers, detailing how to utilise data lakes effectively. It simplifies complex concepts and focuses on practical strategies and the strategic use of data lake technologies like AWS, Azure, and GCP. The book addresses common challenges such as data silos and security. It offers insights into the future of data technologies, empowering managers to harness data for strategic decision-making and innovation.

Link in Amazon Store: https://www.amazon.in/dp/B0D35RCDPD

Recommendation Systems for Managers

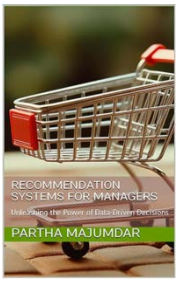

This book demystifies the complexities of data-driven recommendation systems in an easy-to-understand format tailored for managers. This insightful guide traverses Time Series and Market Basket Analysis, AI, ML, and emerging technologies, offering a practical roadmap for implementing these systems. It's an indispensable resource for managers aiming to harness recommendation systems for strategic business decisions in the digital age.

Link in Amazon Store: https://www.amazon.in/dp/B0CXNNSJRC

Learn Emotion Analysis with R

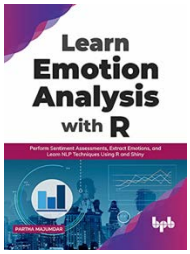

This book is a comprehensive guide to Emotion Analysis using Lexicons, offering a step-by-step code walkthrough for developing Sentiment and Emotion Analysis systems with data from WhatsApp and Twitter. It introduces R and Shiny programming, which is essential for building emotion analysis systems. The discussion then extends to the fundamentals of Sentiment and Emotion Analysis, leading to the creation of Shiny applications tailored for this purpose. The book concludes by developing a specialised tool for analysing emotions from Twitter and WhatsApp data. Additionally, it hints at advancing into Machine Learning for Emotion Analysis, contingent on the availability of labelled data, positioning this as a subsequent step for readers.

Link in Amazon Store: https://www.amazon.com/dp/B096K2SVF2

Starting a New AI Business

This book is a comprehensive guide designed for entrepreneurs looking to harness the power of artificial intelligence to build successful enterprises. Covering everything from defining business purpose and understanding AI fundamentals to exploring innovative business models and identifying market opportunities, this book provides practical insights and strategic guidance. With case studies of industry giants and lessons from ancient wisdom, it equips readers with the tools and knowledge to navigate the AI landscape effectively and achieve sustainable growth.

Link in Amazon Store: https://www.amazon.in/dp/B0CL3YBSF8

Ten Essays on AI

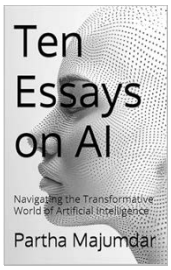

This book offers a deep dive into the world of artificial intelligence, covering critical topics such as data-analytic thinking, data quality, AI strategies, and the evolution of AI technology. This comprehensive collection provides readers with practical insights and theoretical knowledge to harness AI's potential and navigate its challenges effectively.

Link in Amazon Store: https://www.amazon.in/dp/B0D9PMJBWB

Essays on Machine Learning

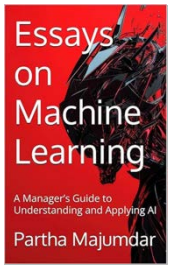

This book is a concise and practical guide for managers and business leaders seeking to understand and apply machine learning in their organisations. It distils complex concepts into clear, accessible insights and covers foundational principles, practical applications, and advanced techniques. Designed as a ready reckoner, it empowers non-technical leaders with the knowledge to make informed decisions in the AI-driven landscape, driving innovation and strategic growth.

Link in Amazon Store: https://www.amazon.in/dp/B0DDHZ78Z7

Navigating Economic Frontiers

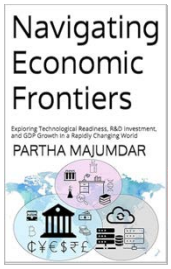

This book dissects the complex interplay between R&D spending, GDP growth, and technological readiness. Framed through an Indian perspective but adaptable for global analysis, this book offers policymakers, business leaders, and data analysts practical insights and Python-based tools to explore these economic dynamics. With real-world correlations and actionable strategies, "Navigating Economic Frontiers" empowers readers to navigate the global economy's evolving landscape, identify growth opportunities, and make data-driven decisions in an interconnected world.

Link in Amazon Store: https://www.amazon.in/dp/B0DLNR2H4Q

Securing Authenticity

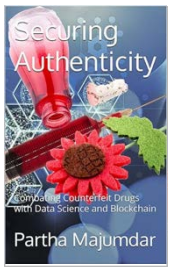

This book leverages data science and blockchain to address the urgent challenge of counterfeit pharmaceuticals. It provides a comprehensive framework for creating a transparent and traceable drug supply chain, ensuring that only genuine medicines reach consumers. By combining innovative technology, strategic partnerships, and robust regulatory compliance, this book empowers stakeholders to collaborate in building safer healthcare systems. Through detailed technical solutions, real-world examples, and practical strategies, "Securing Authenticity" offers a path toward protecting public health and restoring trust in global healthcare markets.

Link in Amazon Store: https://www.amazon.com/dp/B0DMKLQGB6

Linear Programming for Project Management Professionals

This guide provides project management professionals with strategies for project crashing using linear programming, ensuring timely completion and cost efficiency. It introduces basic project management concepts, monitoring techniques, and linear programming problem formulation. The book explains how to solve these problems using Microsoft Excel's Solver and applying time and cost optimisation methods to real-world scenarios. It equips project management teams with a comprehensive toolkit to handle complex challenges effectively.

Link in Amazon Store: https://www.amazon.com/dp/B09PD1GFMY

Gartner Research Analysis

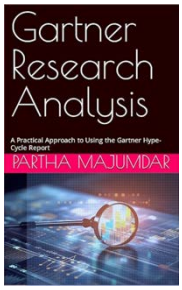

The book provides a clear framework for leveraging insights from the Gartner Hype-Cycle Report, an essential resource for understanding technological trends. It simplifies identifying and evaluating emerging technologies, their developers, and market readiness. A live case study illustrates practical application while emphasising the need for comprehensive research beyond the report. Essential for those seeking strategic technological guidance, this book demystifies the complex data presented in the Gartner Hype Cycle.

Link in Amazon Store: https://www.amazon.com/dp/B0CK582Y2M

Corporate Lessons I Learned

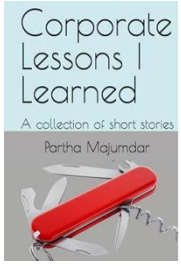

This book encapsulates 34 years of corporate experiences up to 2023, presenting a collection of impactful incidents and interactions that shaped the author's career. Primarily aimed at middle and lower-level managers, it offers humorous and insightful recollections that serve as practical guidelines for navigating daily challenges in the corporate world. The author illustrates valuable lessons learned through various encounters, making it a helpful resource for understanding and excelling in corporate management.

Link in Amazon Store: https://www.amazon.in/dp/B0CL3YBSF8

Rise and Fall

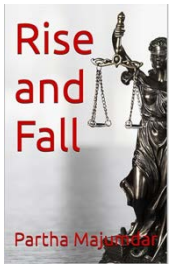

This book follows Mahesh Roy's transformation from a small-town boy to a powerful hotelier entangled in illegal activities. Driven by ambition, Mahesh's world unravels as he confronts the consequences of his choices. In a parallel narrative, lawyer Durga Varma fights to expose the truth behind the mysterious death of her beloved Raj in Mahesh's casino. This gripping novel explores the themes of power, greed, justice, and redemption in a dramatic clash between integrity and corruption.

Link in Amazon Store: https://www.amazon.in/dp/B0DFY5F3VF

Transformation

This book tells the story of Roopesh, a successful yet deeply troubled man who embarks on a journey of change guided by Vedant, a mentor versed in ancient wisdom. Through small, deliberate steps, Roopesh learns to reclaim his life, face his inner demons, and find balance in the everyday. The book explores self-discipline, mindfulness, and the importance of surrendering control, offering a powerful reminder that true transformation begins within and that it's never too late to change.

Link in Amazon Store: https://www.amazon.in/dp/B0DGG5343N

A Journey Back to Paris

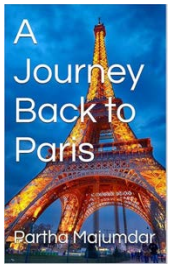

In "A Journey Back to Paris," Paresh and Romila embark on a holiday to Paris that transforms into something more meaningful. Along the way, they encounter Mahua, a distressed young student, sparking a chain of events that redefine their trip. As the trio travels through Paris, Antwerp, and Amsterdam, themes of love, trust, and transformation unfold, making this story a rich blend of personal introspection and cultural adventure. It's a tale of connection, chance, and life's beautiful unpredictability.

Link in Amazon Store: https://www.amazon.in/dp/B0DK2L3XS9

Mutual Fund Investing

This book is a comprehensive guide for middle-class investors in India, simplifying mutual funds. It covers types of mutual funds, differences between open-ended and closed funds, systematic investments, tax implications, and risk assessment. It also teaches advanced techniques like Piotrowski's F-Score and Mohanram's G-Score for building diversified portfolios and evaluating fund performance. Suitable for beginners and seasoned investors, it is essential for achieving financial growth and security through mutual funds.

Link in Amazon Store: https://www.amazon.com/dp/B0CYNG6B12

Creating an Investment Portfolio

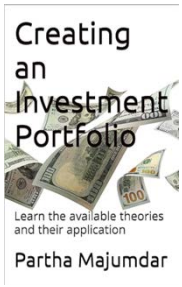

This book delves into the scientific process of making informed investment decisions, highlighting the importance for individuals and corporations. It explores critical theories and applications in portfolio creation, covering various investment vehicles like fixed deposits, mutual funds, and shares, emphasising the necessary mathematics. Additionally, it introduces simple yet widely used tools for investment calculations. Designed to be accessible to a broad audience, this book is an invaluable guide for beginners and experienced investors aiming to enhance their understanding and effectiveness in investing.

Link in Amazon Store: https://www.amazon.com/dp/B0CK99SPKZ

Essay on the Indian Knowledge System – Part 1

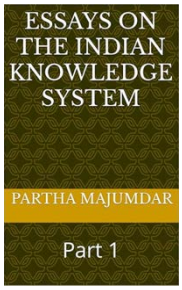

The book delves into the Indian Knowledge System (IKS), a comprehensive approach to compiling, conserving, and disseminating India's rich knowledge heritage across various disciplines such as science, mathematics, social sciences, medicine, philosophy, art, and spirituality. It highlights the global perspective of IKS and its relevance in sharing India's intellectual legacy with the world. The study of Indology, or "Bharatatattva," as it's known in Indian scholarship, further explores the historical, cultural, linguistic, and literary facets of the Indian subcontinent. Through a series of concise essays, this book, one of a trilogy on ancient India, offers insights into Bharatatattva, underscoring India's significant contributions to global knowledge.

Link in Amazon Store: https://www.amazon.com/dp/B0CXNN95TR

Essay on the Indian Knowledge System – Part 2

The book explores ancient India's profound contributions to various scientific disciplines, such as mathematics, astronomy, medicine, and metallurgy. It highlights India's enduring influence on global scientific thought by examining the works of great scholars like Aryabhata and Sushruta. With a unique blend of philosophy and empirical science, this book offers insights into how ancient wisdom continues to shape modern scientific inquiry. It is an essential read for those interested in the history of science.

Link in Amazon Store: https://www.amazon.in/dp/B0DJ7NZH82

Essay on the Indian Knowledge System – Part 3

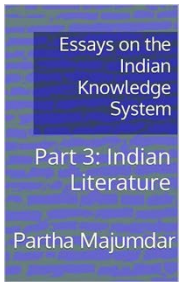

The book comprehensively examines India's literary evolution. From the Vedas and Upanishads to the works of Kalidas, Rabindranath Tagore, and contemporary Indian writers, it covers a wide range of topics. It delves into Indian literature's spiritual and intellectual contributions, exploring its influence on culture, philosophy, and human thought. It is a must-read for scholars, students, and anyone interested in India's profound literary heritage.

Link in Amazon Store: https://www.amazon.in/dp/B0DKDF6CNF

Good People Are Tested the Most

This book explores the lives of seven extraordinary individuals who faced immense challenges and emerged victorious. The book delves into resilience, faith, and the triumph of the human spirit, featuring Bhakt Prahalad, Raja Harishchandra, Lord Shri Ram, Arjun, Hercules, Swami Vivekananda, and Sardar Milkha Singh. Their inspiring stories highlight the importance of maintaining principles, overcoming significant challenges, and the ultimate triumph of good over evil, serving as inspiration for all.

Link in Amazon Store: https://www.amazon.in/dp/B0D1FM44H9

Sailing through the Kali Yug

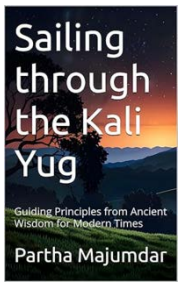

This book explores the relevance of ancient Indian scriptures, especially the Purans, in understanding the complexities of the present age. It introduces the concept of Yugs, highlighting Rishi Krishna Dwaipayan Ved Vyas's role and the Purans' structure. The book details the moral decline of the Kali Yug, starting from Raja Parikshit's reign, and emphasises Dharam's four pillars. It promotes Bhakti and practical spiritual practices as pathways to maintaining integrity and achieving liberation.

Link in Amazon Store: https://www.amazon.in/dp/B0D6M34JCT

The Maha Purans in Brief

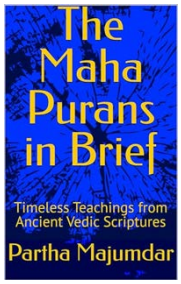

This book distils the essence of eighteen ancient Vedic scriptures, offering insights into their themes, stories, and teachings. The book explores the relevance of Puranic wisdom in modern life, providing guidance on ethics, leadership, environmental sustainability, and personal development. It serves as an accessible guide to understanding the rich spiritual heritage of Indian culture and applying its timeless lessons to contemporary challenges.

Link in Amazon Store: https://www.amazon.in/dp/B0CW8GJ22L

A Collection of Blog Posts

This book is a rich anthology that covers a wide spectrum of topics, including the origins of world religions, the value of hard work, and the importance of strategic leadership. Combining spiritual wisdom, business insights, and practical advice, Majumdar offers readers a thought-provoking blend of reflections designed to inspire personal and professional growth. This collection invites readers to explore ancient lessons while applying them to modern-day challenges in life and work.

Link in Amazon Store: https://www.amazon.com/dp/B0DJHMRMXX

Weekend in Jordan

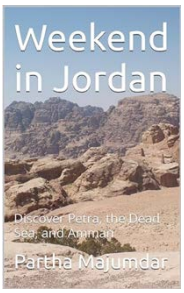

Thanks to the country's visa-on-arrival policy for Indians, the authors spontaneously travelled to Jordan to celebrate their 20th wedding anniversary. Their weekend was filled with memorable experiences, from Petra's historical wonders to the Dead Sea's unique allure and Amman's vibrant city life. Despite its modest size, Jordan's rich offerings left a lasting impression. This book recounts their remarkable journey, providing insights into the treasures of Jordan.

Link in Amazon Store: https://www.amazon.com/dp/B0CK5N6B3W

Elephant Ride in Chang Wangpo

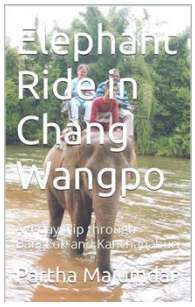

In 2022, Thailand saw a significant influx of approximately 11.5 million tourists, underlining tourism's vital role in its economy, contributing around 6% to the Thai GDP. Reflecting on their past residency in Bangkok from 1996 to 1999, the authors seized a chance to revisit Thailand in 2018, noticing considerable changes. An efficient metro system has alleviated the once notorious Bangkok traffic, enhancing city navigation. While many cherished aspects remained, improvements in the road network and increased attractions enriched their experience. Coinciding with their 26th wedding anniversary, the business trip also included leisure exploration in Bangkok and Kanchanaburi, with a memorable visit to Chang Wangpo, blending nostalgia, discovery, and celebration.

Link in Amazon Store: https://www.amazon.com/dp/B0CKGWH97S

Weekend in South Sikkim

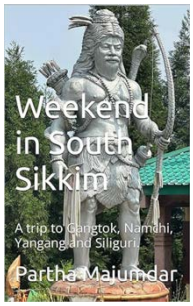

This book explores the less-travelled South Sikkim, diverging from popular tourist spots like Gangtok and Nathu La Pass. It covers captivating destinations such as Tsomgo Lake, Baba Ka Mandir, and Temi Tea Gardens. The authors delve into the cultural and spiritual essence of South Sikkim with visits to Namchi's Char Dham and Samdruptse Monastery. The narrative also extends to Yangang and the Bengal Safari in Siliguri, West Bengal, offering a comprehensive travelogue with diverse experiences.

Link in Amazon Store: https://www.amazon.com/dp/B0CKL1DNTJ

Trips to Dubai

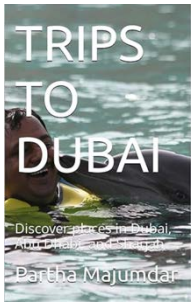

This travelogue unveils the multifaceted allure of Dubai, a top-tier tourist hub known for landmarks like the Burj Khalifa and Burj Al Arab, alongside thrilling experiences such as helicopter rides and dolphin encounters at the Atlantis. It extends beyond Dubai, shedding light on Abu Dhabi and Sharjah attractions, like the adrenaline-pumping Ferrari World and the enchanting Desert Safari. The author shares personal adventures, offering insights into the intricacies of visiting Dubai and navigating the Gulf region, making this book a valuable resource for anyone looking to explore the rich experiences Dubai and its neighbouring emirates offer.

Link in Amazon Store: https://www.amazon.com/dp/B0CKRYQKDN

1-Day Trips from Bengaluru

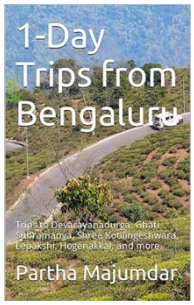

From 1975 to 2023, Bengaluru evolved from a retirees' haven to India's Silicon Valley, also renowned as Garden City. While Bengaluru has numerous tourist attractions and activity hubs, the city's vicinity offers many exploration destinations. This book focuses on day-trip-worthy spots around Bengaluru, places steeped in historical significance. It does not cover prominent cities like Mysuru, Chennai, and Hyderabad, as well as scenic locales like Ooty, Goa, and Kerala, as they need more than a day to tour.

Link in Amazon Store: https://www.amazon.com/dp/B0CLK58KTB

A Trip to the Wagah Border

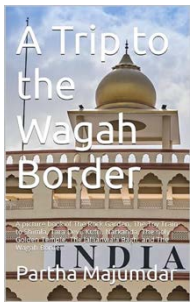

The Wagah Border, straddling India and Pakistan near Amritsar and Lahore, is famed for its ceremonial displays by border forces, symbolising hope amidst strained relations. This checkpoint, pivotal for prisoner exchanges, represents a unique reconciliation potential. On festive occasions, friendly exchanges between the forces foster harmony. The book visually explores Chandigarh, Shimla, Amritsar, and the Wagah Border, highlighting their rich cultures and historical importance.

Link in Amazon Store: https://www.amazon.com/dp/B0CLYTQ6PV

Weekend Getaways from Bengaluru

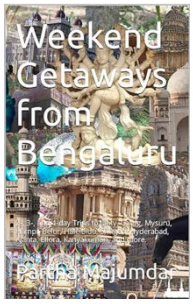

This guidebook enhances the tourism experience in India, emphasising the country's improved accessibility and facilities that cater to all traveller categories. It explicitly outlines short trips from Bengaluru, covering a mix of destinations accessible by road, rail, and air. The book is a resource for planning 2-, 3-, and 4-day excursions to various South Indian locales and select sites in Maharashtra, featuring popular tourist destinations such as Ooty, Kodaikanal, and Mysuru, as well as revered places of worship like Kukke Subramanya and Dharamsthala. It offers practical travel tips, what to anticipate on journeys and insights into each destination's unique offerings.

Link in Amazon Store: https://www.amazon.com/dp/B0CMNRKWQ9

View publication stats